

Web as a Platform



Web History

- Started as a distributed hypermedia platform
 - CERN, Berners-Lee, 1990
- Revolutionised hypermedia
 - Imagine emailing someone a hypermedia deck nowadays!
- Architecture of the Web largely fortuitous
 - W3C and others have since retrofitted/captured the Web's architectural characteristics

The Web broke the rules

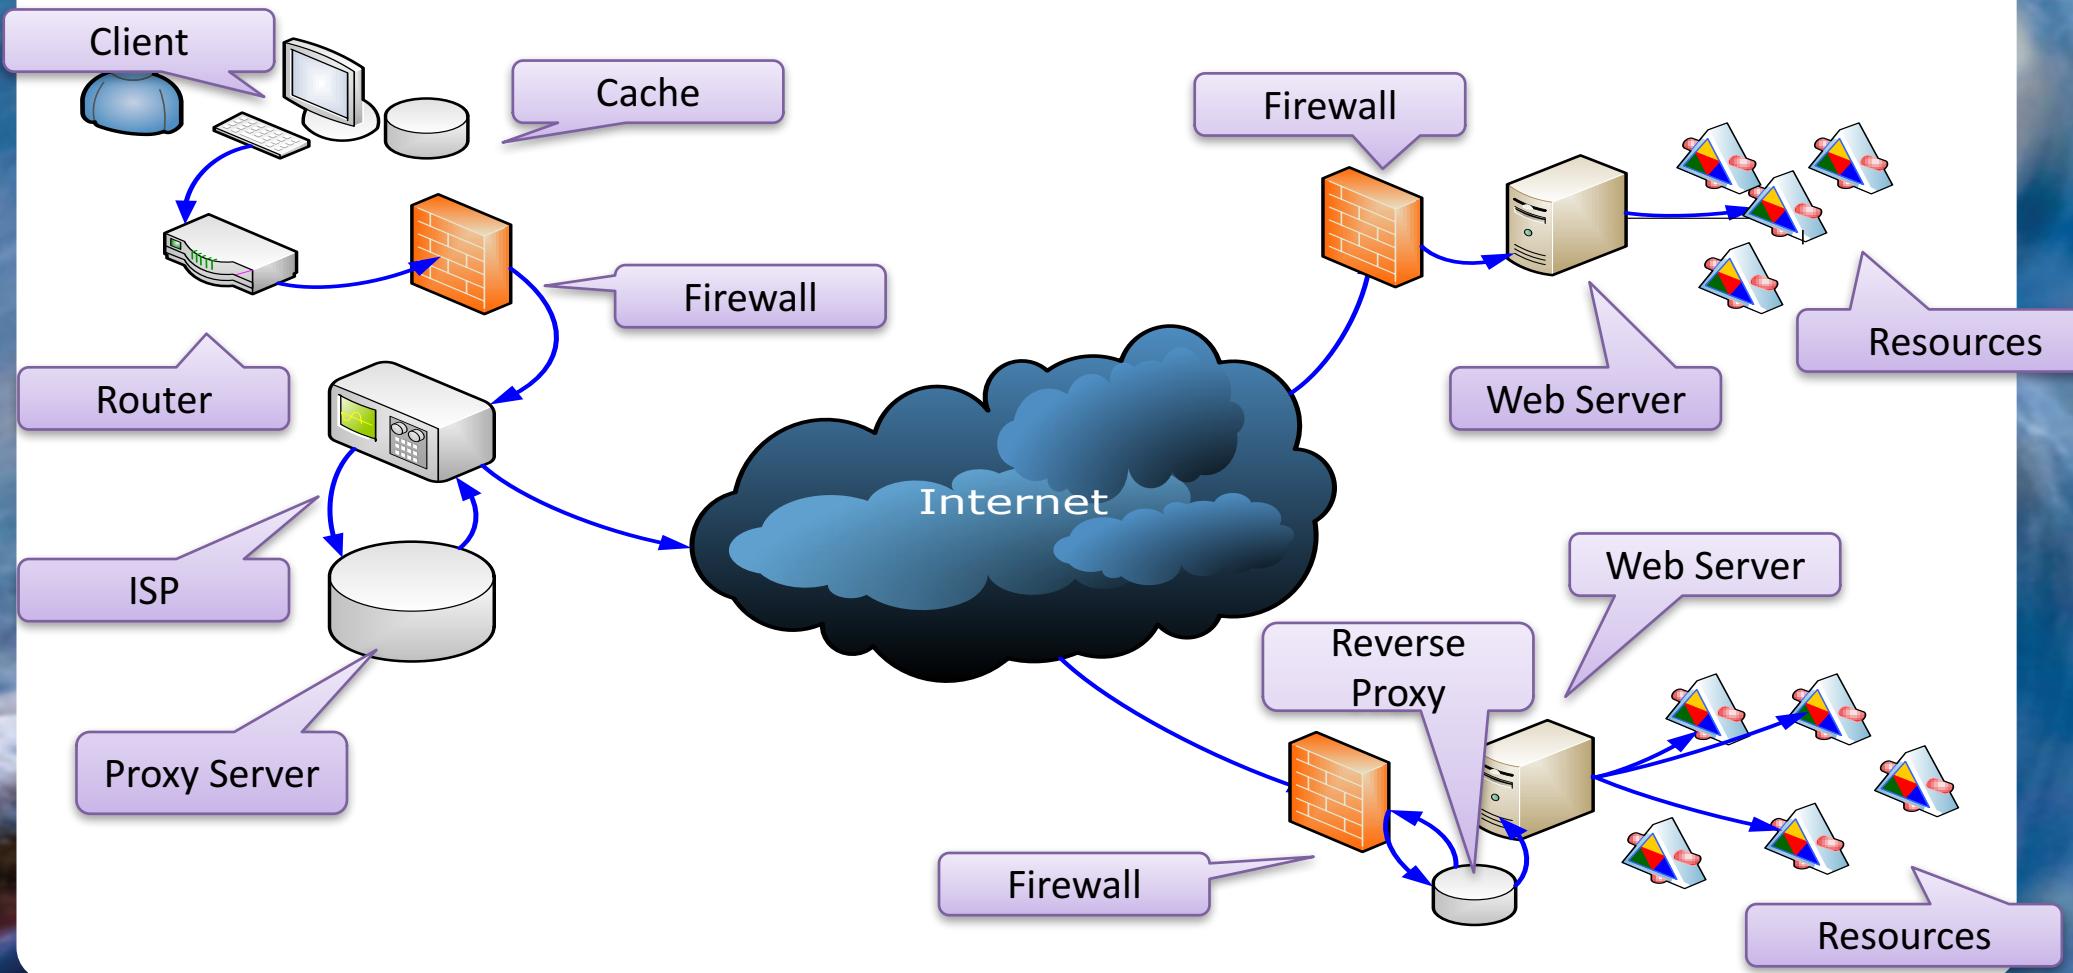


404

Web Fundamentals

- To embrace the Web, we need to understand how it works
- The Web is a distributed hypermedia model
 - It doesn't try to hide that distribution from you!
- Our challenge:
 - Figure out the mapping between our problem domain and the underlying Web platform

Key Actors in the Web Architecture



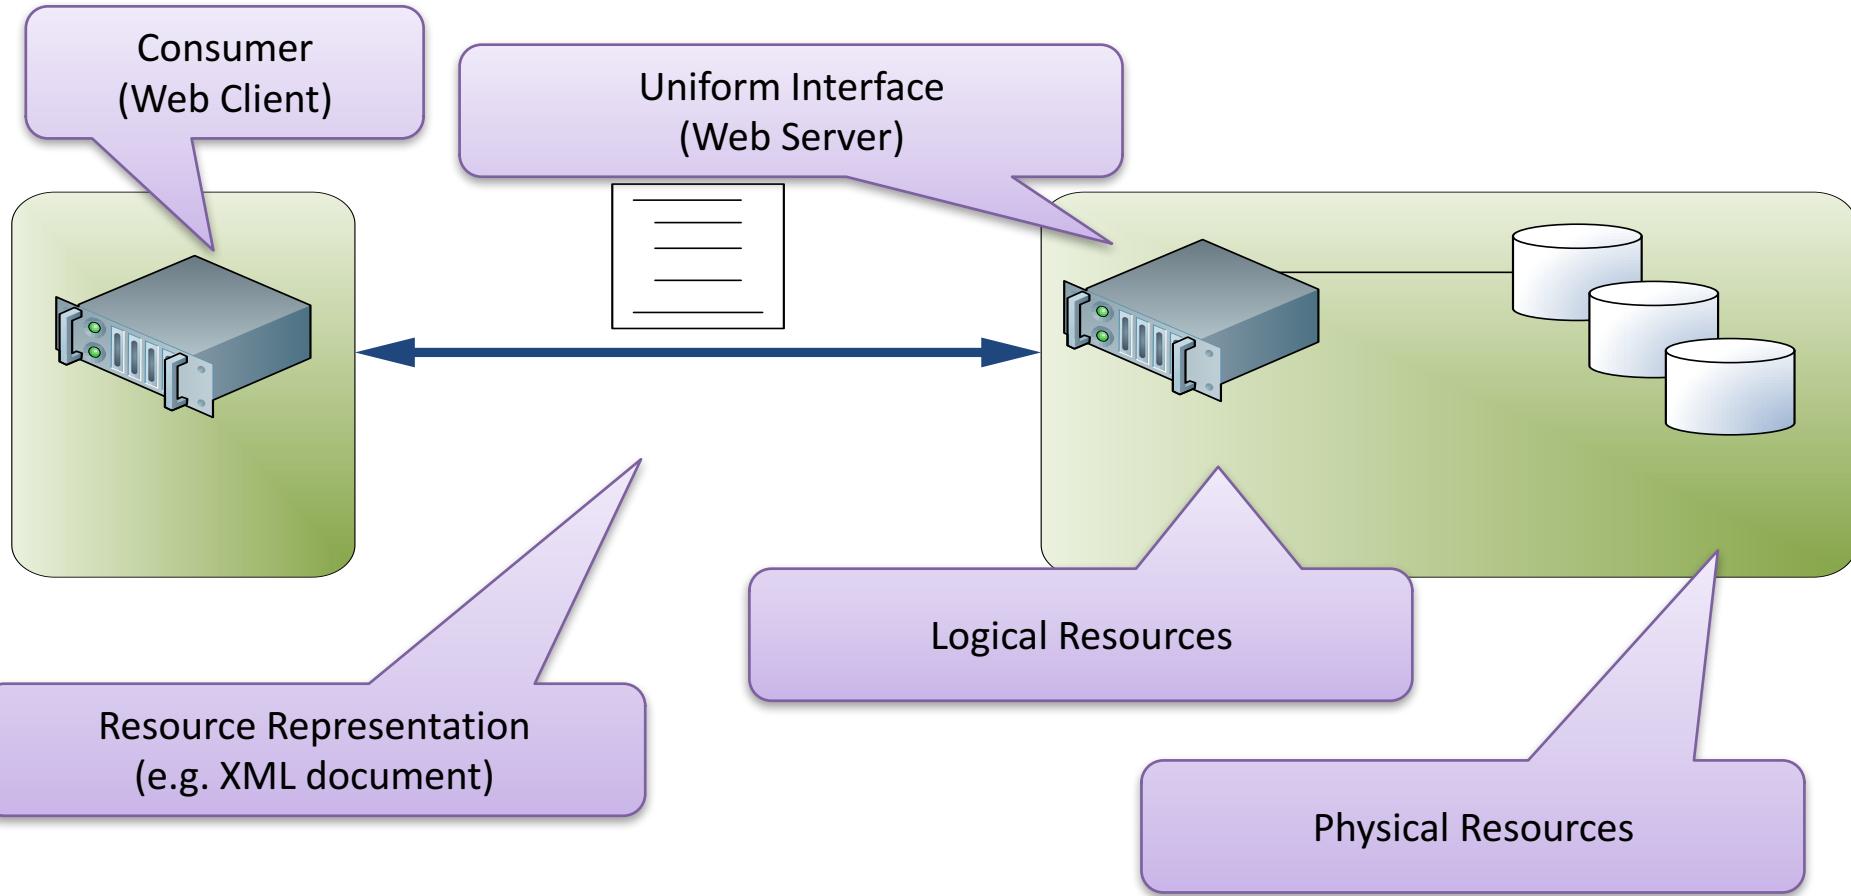
Resources

- A resource is something “interesting” in your system
- Can be anything
 - Spreadsheet (or one of its cells)
 - Blog posting
 - Printer
 - Winning lottery numbers
 - A transaction
 - Others?

Interacting with Resources

- We deal with representations of resources
 - Not the resources themselves
 - “Pass-by-value” semantics
 - Representation can be in any format
 - Any media type
- Each resource implements a standard uniform interface
 - Typically the HTTP interface
- Resources have names and addresses (URIs)
 - Typically HTTP URIs (aka URLs)

Resource Architecture



Resource Representations

- Making your system Web-friendly increases its surface area
 - You expose many resources, rather than fewer endpoints
- Each resource has one or more representations
 - Representations like JSON or XML are good for the programmatic Web
- Moving representations across the network is the way we transact work in a Web-native system

URIs

Part

Identity

Addressing

Description

Distinguishes one entity from another

Means to interact/communicate with entity

Identity

An identifier
not
The identifier

Bounded in time
Cool URIs don't change
URIs change
Get over it

```
<entry>
  <id>urn:uuid:b436fda6-93f5-4c00-98a3-06b62c3d31b8</id>
  <title type="text">promotion cancelled</title>
  <updated>2008-09-10T13:57:14+01:00</updated>
  <link rel="self"
    href="http://restbucks.com/notifications/b436fda6"/>
  <content type="application/xml">
    ...
  </content>
</entry>
```

Addressing: portable, portmanteau, declarative structure

<protocol>://<host>:<port>/<path>

http://restbucks.com:8080/order/123

Vina +
Host

Protocol

Determines how rest of URI is to be interpreted

Host

HTTP library
User agent
Cache
Server

Host

Identifies host domain name

DNS server
Maps host domain name to IP address

Path

(Logical) path to resource

Host-specific
Resolves path, extracts parameters

Scalability

- Web is truly Internet-scale
 - Loose coupling
 - Growth of the Web in one place is not impacted by changes in other places
 - Uniform interface
 - HTTP defines a standard interface for all actors on the Web
 - Replication and caching is baked into this model
 - Caches have the same interface as real resources!
 - Stateless model
 - Supports horizontal scaling

Fault Tolerant

- The Web is stateless
 - All information required to process a request must be present in that request
 - Sessions are still plausible, but must be handled in a Web-consistent manner
 - Modelled as resources!
- Statelessness means easy replication
 - One Web server is replaceable with another
 - Easy fail-over, horizontal scaling

Recoverable

- The Web places emphasis on repeatable information retrieval
 - GET is idempotent
 - Library of Congress found this the hard way!
 - In failure cases, can safely repeat GET on resources
- HTTP verbs plus rich error handling help to remove guesswork from recovery
 - HTTP statuses tell you what happened!
 - Some verbs (e.g. PUT, DELETE) are safe to repeat

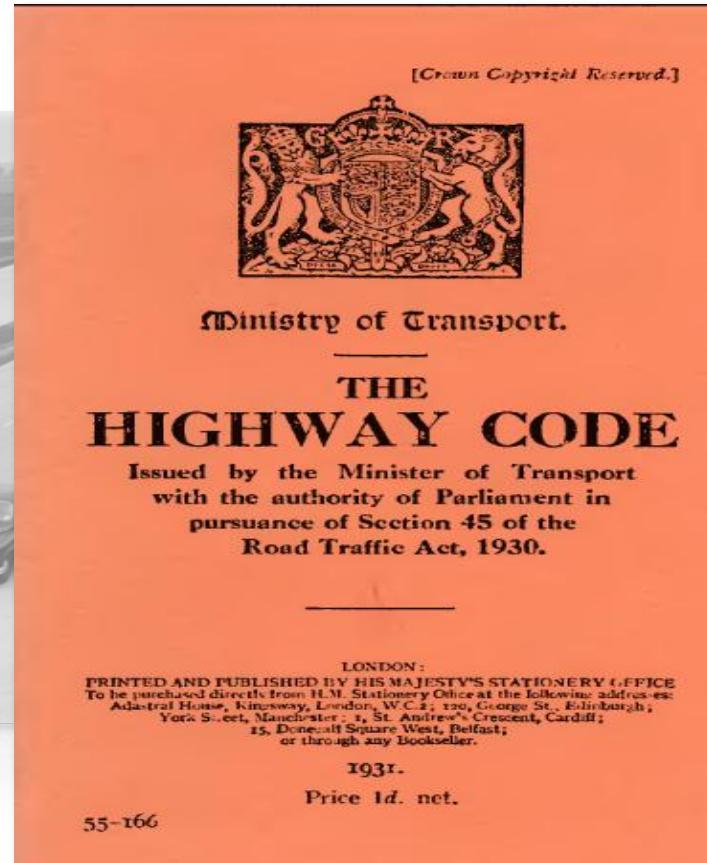
Secure

- HTTPS is a mature technology
 - Based on SSL for secure point-to-point information retrieval
- Isn't sympathetic to Web architecture
 - Restricted caching opportunities
- Higher-order protocols like Atom are starting to change this...
 - Encrypt parts of a resource representation, not the transport channel
 - OK to cache!

Loosely Coupled

- Adding a Web site to the WWW does not affect any other existing sites
- All Web actors support the same, uniform interface
 - Easy to plumb new actors into the big wide web
 - Caches, proxies, servers, resources, etc

HTTP = ???



HTTP Fundamentals



HTTP: an application transfer protocol

Term	Description
Application	Transacted at application layer of network stack, in user agent and server software
Transfer	For the purpose of coordinating the transfer of documents
Protocol	According to rules that determine legitimate interactions between client and server

The HTTP Verbs

- Retrieve a representation of a resource: **GET**
- Create a new resource: **PUT** to a new URI, or **POST** to an existing URI
- Modify an existing resource: **PUT** to an existing URI
- Delete an existing resource: **DELETE**
- Get metadata about an existing resource: **HEAD**
- See which of the verbs the resource understands: **OPTIONS**

Decreasing likelihood of being understood by a Web server today

HEAD Semantics

- HEAD is like GET, except it only retrieves metadata
- Request

```
HEAD /order/1234 HTTP 1.1
```

```
Host: restbucks.com
```

- Response

```
200 OK
```

```
Content-Type:
```

```
application/vnd.restbucks+xml
```

```
Last-Modified: 2023-05-05T00:34Z
```

```
Etag: 77fe653b
```

Useful for caching,
performance

OPTIONS Semantics

- Asks which methods are supported by a resource
 - Easy to spot read-only resources for example
- Request

OPTIONS /o 1.0.2.1 HTTP/1.1

Host: rest

You can only read and add to
this resource, may change
over time

- Response

200 OK

Allow: GET, HEAD, POST

HTTP Status Codes

- The HTTP status codes provide metadata about the state of resources
- They are part of what makes the Web a rich platform for building distributed systems
- They cover five broad categories
 - 1xx - Metadata
 - 2xx – Everything's fine
 - 3xx – Redirection
 - 4xx – Client did something wrong
 - 5xx – Server did a bad thing
- There are a handful of these codes that we need to know in more detail

1XX

- 100 – Continue
 - The operation will be accepted by the service
 - The “look before you leap” pattern
 - Use with the Expect header

- Request

POST /orders HTTP/1.1

Host: restbucks.com

Content-Type: application/vnd.restbucks+xml

Content-Length: 10216

Expect: 100-continue

- Response

HTTP/1.1 100 Continue

OR

HTTP/1.1 417 Expectation Failed

2xx

- 200 – OK
 - The server successfully completed whatever the client asked of it
- 201 – Created
 - Sent when a new resource is created at the client's request via POST
 - Location header should contain the URI to the newly created resource
- 202 – Accepted
 - Client's request can't be handled in a timely manner
 - Location header should contain a URI to the resource that will eventually be exposed to fulfil the client's expectations

More 2xx Codes

- 203 – Non-Authoritative Information
 - Much like 200, except the client knows not to place full trust in any headers since they could have come from 3rd parties or be cached etc.
- 204 – No Content
 - The server declines to send back a representation
 - Perhaps because the associated resource doesn't have one
 - Used like an “ack”
 - Prominent in AJAX applications
- 206 – Partial Content
 - Optimisation used in failure cases to support partial GETs
 - Request Content-Range header must specify the byte range of the resource representation it wants
 - Response headers must contain Date;
 - ETag and Content-Location headers must be consistent with the original request not the current values

3xx

- 301 – Multiple Choices
 - Response Location header should contain the preferred URI
 - Message body can contain list of URIs
 - In XHTML possibly
 - In general avoid being ambiguous!
- 302 – Moved Permanently
 - Location header contains the new location of the resource
- 303 – See Other
 - Location header contains the location of an alternative resource
 - Used for redirection

More 3xx

- 304 – Not Modified
 - The resource hasn't changed, use the existing representation
 - Used in conjunction with conditional GET
 - Client sends the `If-Modified-Since` header
 - Response `Date` header must be set
 - Response `Etag` and `Content-Location` headers must be same as original representation
- 307 – Temporary Redirect
 - The request hasn't been processed, because the resource has moved
 - Client must resubmit request to the URI in the response `Location` header

4xx

- 400 – Bad Request
 - The client has PUT or POST a resource representation that is in the right format, but contains invalid information
- 401 – Unauthorized
 - Proper credentials to operate on a resource weren't provided
 - Response `WWW-Authenticate` header contains the type of authentication the server expects
 - Basic, digest, WSSE, etc
 - Don't leak information!
 - Consider 404 in these situations

More 4xx

- 403 – Forbidden
 - The client request is OK, but the server doesn't want to process it
 - E.g. Restricted by IP address
 - Implies that resource exists, beware leaking information
- 404 – Not Found
 - The standard catch-all response
 - May be a lie to prevent 401 or 403 information leakage

Even more 4xx

- 405 – Method Not Allowed
 - The resource doesn't support a given method
 - The response Allow header lists the verbs the resource understands
 - E.g. Allow: GET, POST, PUT
- 406 – Not Acceptable
 - The client places too many restrictions on the resource representation via the Accept-* header in the request
 - The server can't satisfy any of those representations

Yet More 4xx

- 409 – Conflict
 - Tried to change the state of the resource to something the server won't allow
 - E.g. Trying to DELETE something that doesn't exist
- 410 – Gone
 - The resource has gone, permanently.
 - Don't send in response to DELETE
 - The client won't know if it was deleted, or if it was gone and the delete failed
- 411 – Length Required
 - If a request (POST, PUT) contains a representation, it should set the Content-Length header
 - The server *might* demand this, and interrupt the client request

Still more 4xx

- 412 – Precondition Failed
 - Server/resource couldn't meet one or more preconditions
 - As specified in the request header
 - E.g. Using `If-Unmodified-Since` and `PUT` to modify a resource provided it hasn't been changed by others
- 413 – Request Entity Too Large
 - Response comes with the `Retry-After` header in the hope that the failure is transient

Final 4xx Codes

- 414 – Request URI Too Long
 - You've got a hopeless HTTP server, or ridiculous URI names!
- 415 Unsupported Media Type
 - E.g. If the server resource expects JSON but the client sends XML

5xx Codes

- 500 – Internal Server Error
 - The normal response when we're lazy ☺
- 501 – Not Implemented
 - The client tried to use something in HTTP which this server doesn't support
 - Keep HTTP use simple, get a better HTTP server
- 502 – Bad Gateway
 - A proxy failed
 - Doesn't help us much ☹

More 5xx Codes

- 503 – Service Unavailable
 - The HTTP server is up, but not supporting resource communication properly
 - Server may send a `Retry-After` header, assuming the fault is transient

HTTP Headers

- Headers provide metadata to assist processing
 - Identify resource representation format (media type), length of payload, supported verbs, etc
- HTTP defines a wealth of these
 - And like status codes they are our building blocks for robust service implementations

Must-know Headers

- Authorization
 - Contains credentials (basic, digest, WSSE, etc)
 - Extensible
- Content-Length
 - Length of payload, in bytes
- Content-Type
 - The resource representation form
 - E.g. application/vnd.restbucks+xml,
application/xhtml+xml

More Must-Know Headers

- **Etag:** Etag : "bne3t6s"
 - Opaque identifier – think “checksum” for resource representations
- **If-None-Match**
 - Execute conditional operation if resource **has** changed
 - Conditional GET: 304 Not Modified
 - Other conditional operations: 412 Precondition Failed
- **If-Match**
 - Mainly other conditional operations (PUT, POST, DELETE)
 - Execute conditional operation if resource **hasn't** changed
 - 412 Precondition Failed

Some More Must-Know Headers

- **Last-Modified:** Last-modified: Wed, 03 Mar 2010 20:58:03 GMT
 - Timestamp
- **If-Modified-Since**
 - Execute conditional operation if resource *has* changed
 - Conditional GET: 304 Not Modified
 - Other conditional operations: 412 Precondition Failed
- **If-Not-Modified-Since**
 - Mainly other conditional operations (PUT, POST, DELETE)
 - Execute conditional operation if resource *hasn't* changed
 - 412 Precondition Failed

Yet More Must-Know Headers

- Location
 - Used to flag the location of a created/moved resource
 - In combination with:
 - 201 Created, 301 Moved Permanently, 302 Found, 307 Temporary Redirect, 300 Multiple Choices, 303 See Other
- User-Agent
 - Tells the server side what the client-side capabilities are
 - Should not be used in the programmable Web!

Final Must-Know Headers

- WWW-Authenticate
 - Used with 401 status
 - Informs client what authentication is needed
- Date
 - Mandatory!
 - Timestamps on request and response
- Host
 - Contains the domain-name part of the URI

Useful Headers

- Accept
 - Client tells server what formats it wants
 - Can externalise this in URI names in the general case
- Accept-Encoding
 - Client tells server that it can compress data to save bandwidth
 - Client specifies the compression algorithm it understands
- Content-Encoding
 - Server-side partner of Accept-Encoding

More Useful Headers

- Allow
 - Server tells client what verbs are allowed for the requested resource (used in combination with OPTIONS)
- Cache-Control
 - Metadata for caches, tells them how to cache (or not) the resource representation
 - And for how long etc.
- Content-MD5
 - Cryptographic checksum of body
 - Useful integrity check, has computation cost

Yet More Useful Headers

- Expect
 - A conditional – client asks if it's OK to proceed by expecting 100-Continue
 - Server either responds with 100 or 417 – Expectation Failed
- Expires
 - Server tells client or proxy server that representation can be safely cached until a certain time
- If-Match
 - Used for ETag comparison
 - Opposite of If-None-Match

Final Useful Headers

- If-Unmodified-Since
 - Useful for conditional PUT/POST
 - Make sure the resource hasn't changed while you're been manipulating it
 - Compare with If-Modified-Since
- Range
 - Specify part of a resource representation (a byte range) that you need – aka partial GET
 - Useful for failure/recovery scenarios

Less Often-Used Headers

- Retry-After
 - Resource or server is out of action for the specified time
 - Usually associated with 413 – Request Entity Too Large, or one of the 5xx server error statuses
- Content-Location
 - Header gives the canonical URI of the resource
 - Client might be using a different URI to access that resource

HTTP RCF 2616 is Authoritative

- The statuses and headers here are a sample of the full range of headers in the HTTP spec
- They spec contains more than we discuss here
- It is authoritative about usage
- And it's a good thing to keep handy when you're working on a Web-based distributed system!

The Platform is the Web



Web as a Platform

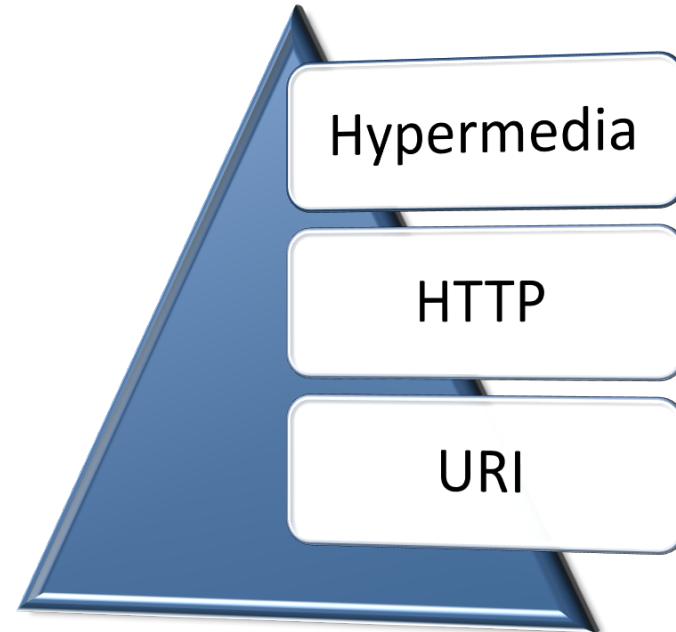
- The Web is a reliable integration platform
 - Not just for human content
- HTTP is a very well-thought-out protocol for manipulating objects at a distance
 - Concurrency control, caching, fault tolerance, resume, coordination all baked in
- This session takes a deep look at what's on offer

The Web is middleware

- The Web is a middleware platform which is...
 - Globally deployed
 - Has reach
 - Is mature
 - And is a reality in every part of our lives
- Which makes it interesting for distributed systems like microservices

Leonard Richardson's Web service maturity heuristic

- REST is cool
- But the Web is permissive
- There's a spectrum of maturity of service styles
 - From completely bonkers to completely RESTful
- We'll use the Richardson maturity model to frame these kinds of discussions
 - Level 0 to Level 3
 - Web-ignorant to RESTful!



Caches
Microformats
CDNs
Media Types
HTTP
OAuth
SemWeb
URIs

REST is an Architectural style suited to the Web...

...but it is not mandatory

Why the Web? Why be RESTful?

- Scalable
- Fault-tolerant
- Recoverable
- Secure
- Loosely coupled
- Precisely the same characteristics we want in business software systems!

REST isn't...

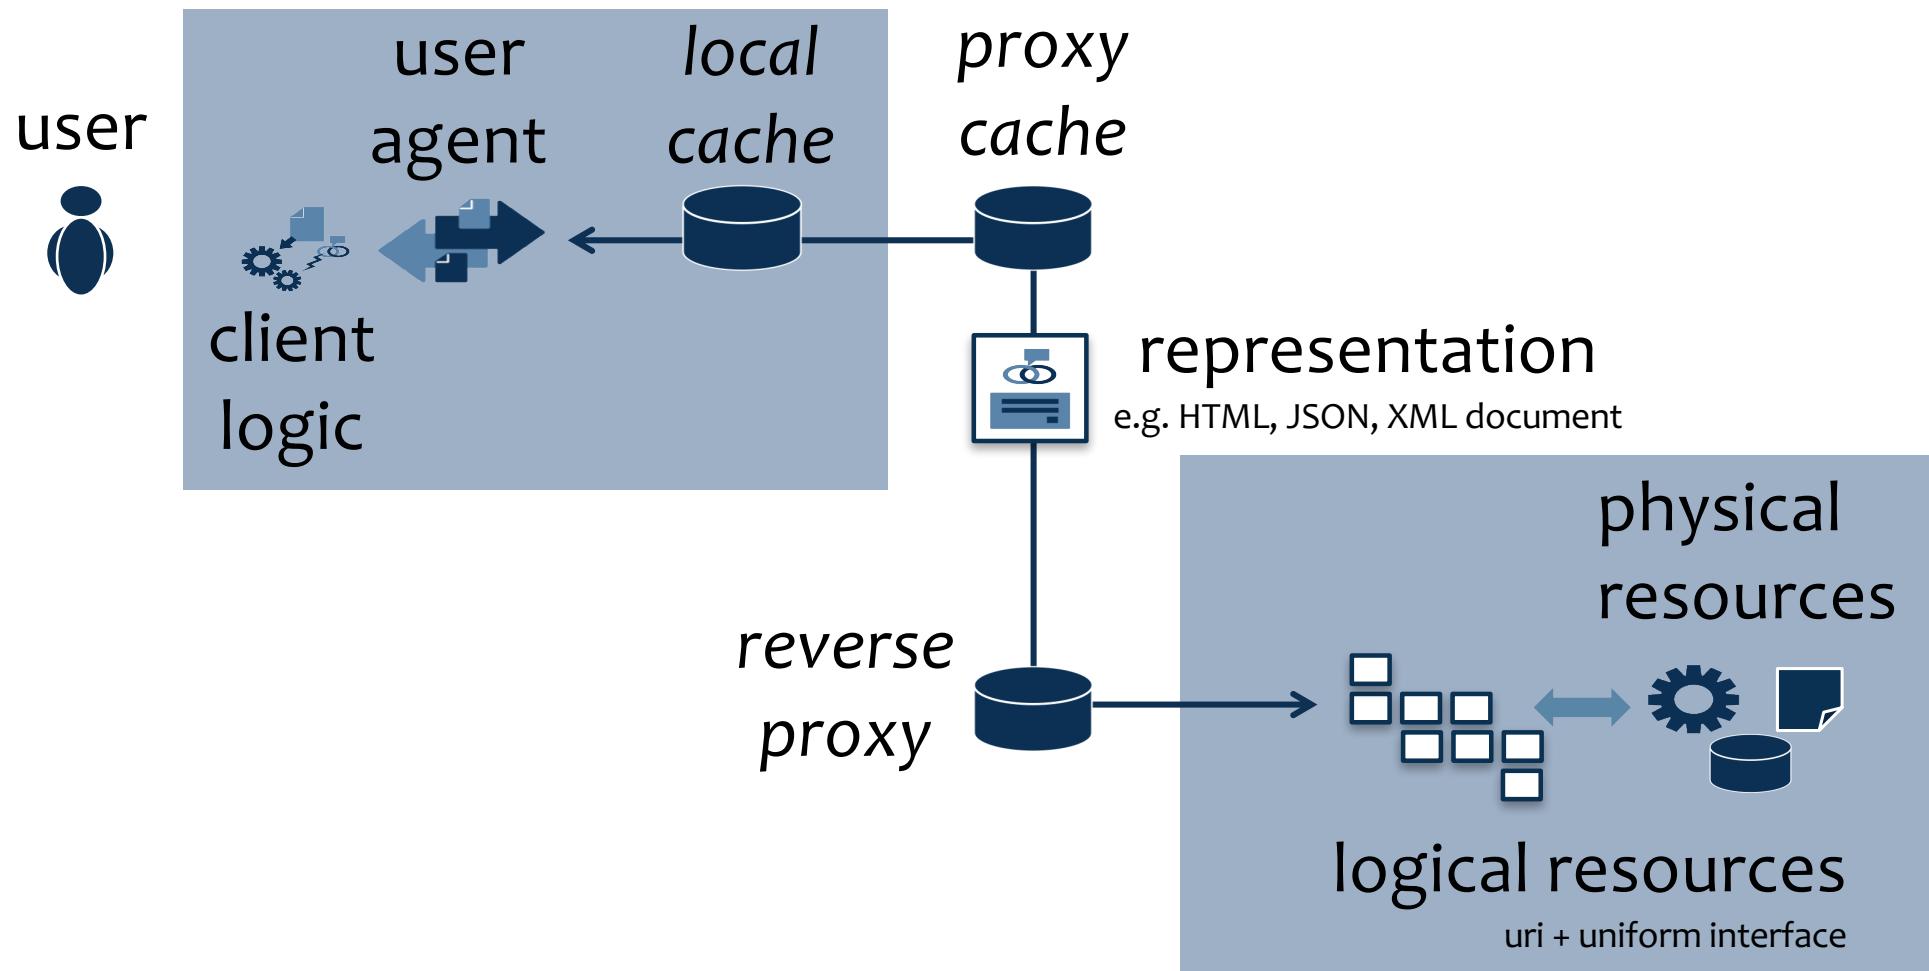
- Just pretty URIs
- Just XML or JSON
- URI templates
- AJAX applications

But you could be forgiven for thinking it is!

Web Architecture



Web Architecture



REST: A Formal Description of the Architecture of the Web

- Roy Fielding, Chapter 5, “Architectural Styles and the Design of Network-based Software Architectures”
 - <https://www.ics.uci.edu/~fielding/pubs/dissertation/top.htm>
- **Representational State Transfer**
 - Transfer representations of state (documents) between client and server in order to progress the state of an application
- An architectural **style**
 - Application of design constraints
 - For network-based applications
 - To induce the architectural properties of the Web

REST's Architectural Constraints (1)

Constraint	Induces	Tradeoff
Client-Server	Separation of concerns UI portability Simplified server Multiple organizational domains	
Stateless	Visibility Reliability Scalability	Decreased network performance Server does not control application state
Caching	Reduced latency Efficiency Scalability	Stale data reduces <i>reliability</i>

REST's Architectural Constraints (2)

Constraint	Induces	Tradeoff
Uniform Interface	Visibility Decoupled implementation Independent evolution	Standardized rather than application-specific format degrades efficiency
Layered System	Reduced complexity Substrate independence Encapsulate legacy components Shared intermediaries Load balancing	Increased latency
(Optional) Code-on-Demand	Simplified clients Extensibility	Reduced visibility

REST's Uniform Interface Constraints

Constraint	Description
Identification of resources	Each important resource is uniquely <i>identifiable</i> One resource identifier mechanism: e.g. URI
Manipulation of resources through representations	Request – via document exchange – rather than tell
Uniform access semantics	Uniform document exchange protocol: e.g. HTTP
Self-descriptive messages	Intention “written on the wire” e.g. HTTP + media-type processing model
Hypermedia as the engine of application state	Hypermedia systems transform application state

<http://roy.gbiv.com/untangled/2008/rest-apis-must-be-hypertext-driven>

What is a Resource?

- Address + identity (URI)
- Encapsulates state (data)
- Manipulated through representations (e.g. HTML, JSON, XML)
- Supports HTTP uniform interface

<http://restbucks.com/quotes/1234>



```
<shop xmlns="http://schemas.restbucks.com/shop"
      xmlns:rb="http://relations.restbucks.com/">
  <items>
    <item>
      <description>Costa Rica Tarrazu</description>
      <amount>250g</amount>
      <price currency="GBP">4.40</price>
    </item>
  </items>
  <link rel="rb:order-form"
        href="http://restbucks.com/order-forms/1234"/>
</shop>
```

GET
PUT
POST
DELETE



Web Technologies

- URIs
- Media Types
- HTTP

URIs are Names for Resources

`http://restbucks.com:8080/order/123`

Identity

Differentiates one resource from another

Addressing

Means for accessing and manipulating resource representations

Each resource has **at least one** address

Every resource implements the same uniform interface – HTTP

Not Semantics

```
<link href="http://restbucks.com:8080/order/123" rel="order"/>
```

URI – Declarative Structure Targeting Different Parts of Network

<protocol>://<host>:<port>/<path>

http://restbucks.com:8080/order/123

What	Protocol	Host	Path
How	Determines how rest of URI is to be interpreted User agent Client software	Identifies host domain name DNS server Maps host domain name to IP address	(Logical) path to resource Host-specific Resolves path, extracts parameters

Media Types

- Used for labelling and formatting resource representations
- Some types are special because they work in harmony with the Web
 - We call these “hypermedia formats”
- A media type defines a processing model
 - If the client understands the type, it can process the content
- Processing model may include:
 - Representation format (may include schemas)
 - Supported operations (methods, headers and status codes)
 - Definitions of hypermedia controls (links, forms)
 - Application semantics (link relations)

Hypermedia Formats

- **JSON Hypertext Application Language**
 - Supports links and embedded resources
 - No intrinsic support for forms/operations
- **JSON for Linked Documents**
 - Can be used to enhance existing JSON
 - Use (unofficial) HYDRA vocabulary for operations
- **Collection+JSON**
 - Query templates
 - semantic profiles (application semantics for message body)
- **HTML**
 - Intrinsic support for links and forms
 - Application semantics using e.g. microformats
- **Atom**
 - Atom Syndication Format for structure
 - Atom Publication Protocol for operations
 - Supports nested representations using other media types

JSON

```
{  
  "playerId": "1234567890",  
  "alias": "soofaloofa",  
  "displayName": "Kevin Sookocheff",  
  "profilePhotoUrl": "https://api.example.com/player/1234567890/avatar.png"  
}
```

<http://sookocheff.com/post/api/on-choosing-a-hypermedia-format/>

JSON-LD

```
{  
  "@context": {  
    "name": "https://schema.org/name",  
    "alternateName": "https://schema.org/alternateName",  
    "image": {  
      "@id": "https://schema.org/image",  
      "@type": "@id"  
    },  
    "friends": {  
      "@container": "@set"  
    }  
  },  
  "@id": "https://api.example.com/player/1234567890",  
  "playerId": "1234567890",  
  "name": "Kevin Sookocheff",  
  "alternateName": "soofaloofa",  
  "image": "https://api.example.com/player/1234567890/avatar.png",  
  "friends": [  
    {  
      "@id": "https://api.example.com/player/1895638109"  
    },  
    {  
      "@id": "https://api.example.com/player/8371023509"  
    }  
  ]  
}
```

<http://sookocheff.com/post/api/on-choosing-a-hypermedia-format/>

HAL

```
{  
    "_links": {  
        "self": { "href": "https://api.example.com/player/1234567890" },  
        "https://api.example.com/docs/rels/friends": {  
            "href": "https://api.example.com/player/1234567890/friends" }  
    },  
    "playerId": "1234567890",  
    "name": "Kevin Sookocheff",  
    "alternateName": "soofaloofa",  
    "image": "https://api.example.com/player/1234567890/avatar.png"  
}
```

<http://sookocheff.com/post/api/on-choosing-a-hypermedia-format/>

Collection+JSON

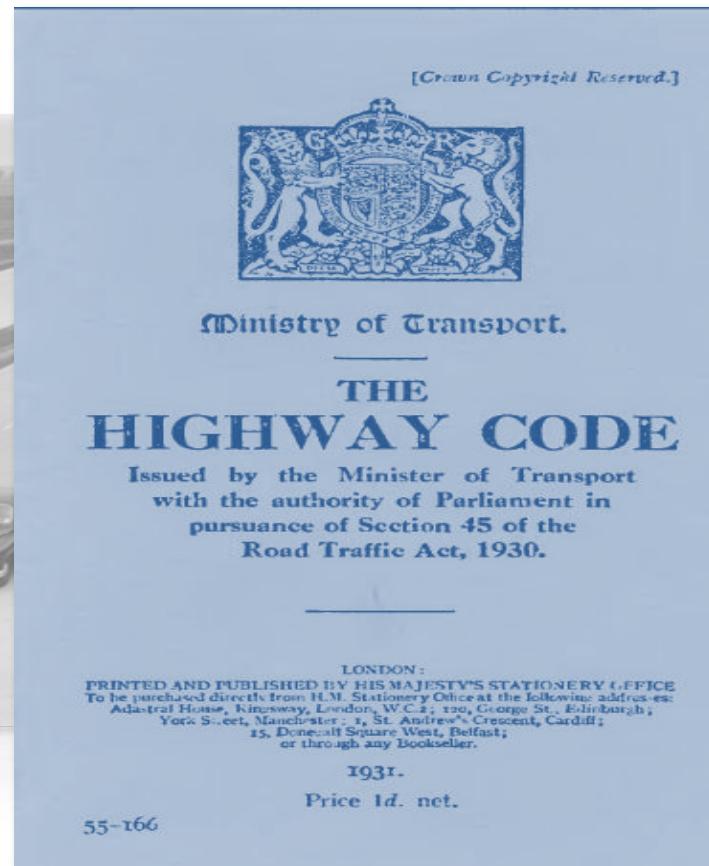
```
{  
  "collection": {  
    "version": "1.0",  
    "href": "https://api.example.com/player",  
    "items": [  
      {  
        "href": "https://api.example.com/player/1234567890",  
        "data": [  
          {"name": "playerId", "value": "1234567890", "prompt": "Identifier"},  
          {"name": "name", "value": "Kevin Sookocheff", "prompt": "Full Name"},  
          {"name": "alternateName", "value": "soofaloofa", "prompt": "Alias"}  
        ],  
        "links": [  
          {"rel": "image",  
           "href": "https://api.example.com/player/1234567890/avatar.png",  
           "prompt": "Avatar",  
           "render": "image" },  
          {"rel": "friends",  
           "href": "https://api.example.com/player/1234567890/friends",  
           "prompt": "Friends" }  
        ]  
      }  
    ]  
  }  
}
```

<http://sookocheff.com/post/api/on-choosing-a-hypermedia-format/>

Decision Criteria

- Plain text
 - Ideal for very simple data payloads
- HAL, JSON-LD, Collection+JSON
 - Best choice for JavaScript (e.g. browser-based) clients
 - HAL & JSON-LD preserve domain/document structure; Collection+JSON is message-oriented
 - <http://sookacheff.com/post/api/on-choosing-a-hypermedia-format/>
- HTML
 - Links and forms
 - Viewable in the browser
- Atom
 - Nested media types
 - Viewable in a feed reader
- Remember, you can mix-and-match!
 - Use content-negotiation to service different clients
 - E.g. plain text, HTML and Atom

HTTP is a (Document Transfer) Application Protocol



HTTP's Uniform Interface

Every resource accessed and manipulated in the same way using:

Status codes

Coordination

Headers

Message processing context

Availability and consistency

Methods

Reliability

Status Codes Coordinate Exchanges

HTTP/1.1 200 OK

HTTP/1.1 201 Created

Location: <http://restbucks.com/quotes/1234>

HTTP/1.1 202 Accepted

Location: <http://restbucks.com/orders/9876>

HTTP/1.1 303 See Other

Location: <http://restbucks.com/orders/9876>

HTTP/1.1 401 Unauthorized

WWW-Authenticate: Basic realm="restbucks"

Even 404 Provides Useful Coordination Information

Client logic (pseudocode)

```
if (GET(/orders/1234) == 200 &&
    GET(/orders/1234/payment) == 200 &&
    GET(/orders/1234/cancellation) == 404 &&
    OPTIONS(/orders/1234/payment).Contains(POST)) {
    applicationState = AwaitingPayment
}

if (POST(/orders/1234/payment) == 405 &&
    GET(/orders/1234/cancellation) == 404) {
    applicationState = Paid
}
```

HTTP Headers

Parameterize the request

Host

Information about the client

Accept

Message metadata

Cache-Control, Date

Information about the server

Server

Response generation

Vary

Further access to the resource

ETag

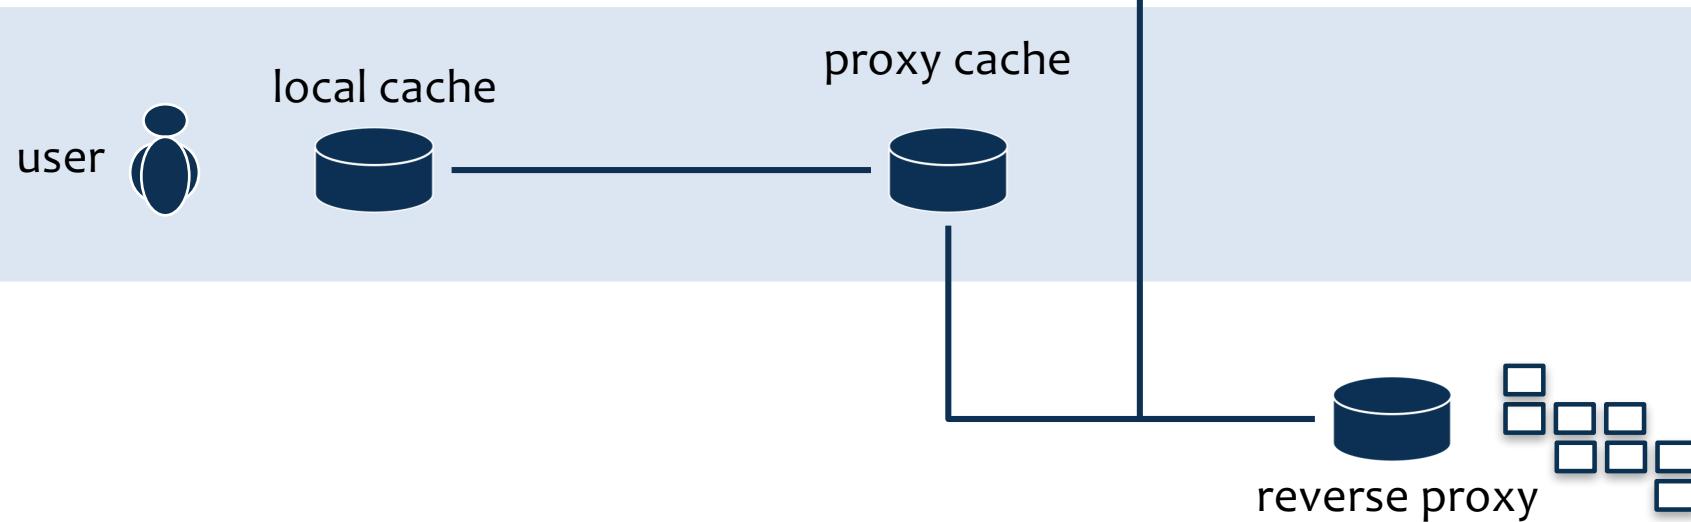
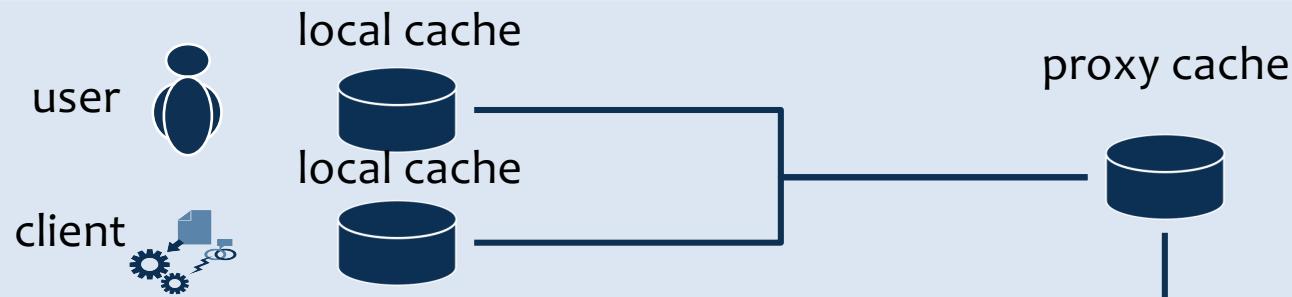
Entity body metadata

Content-Type

Resource metadata

Last-Modified

Caching for Availability and Federated Load



Cache-Control: public, max-age=86400

ETags for Consistency

ETag

Opaque “checksum” of resource state

When resource changes, entity tag changes

Conditional GET

If-None-Match: <etag value>

“Return a full response only if the resource *has changed*”

304 Not Modified

Conditional POST/PUT

If-Match: <etag value>

“Process this request only if the resource *hasn’t changed*”

412 Precondition Failed

Etag Example – First Request

Request

```
GET /orders/9876 HTTP/1.1  
Host: restbucks.com
```

Response

```
HTTP/1.1 200 OK  
ETag: "4d3e88c9"  
Content-Type: application/restbucks+xml
```

```
<shop xmlns="http://schemas.restbucks.com/shop"  
      xmlns:rb="http://relations.restbucks.com/">  
    <status>Awaiting Payment</status>  
</shop>
```

Etag Example – Second Request

Request

```
GET /orders/9876 HTTP/1.1  
Host: restbucks.com  
If-None-Match: "4d3e88c9"
```

Response

```
HTTP/1.1 304 Not Modified
```

Etag Example – Third Request

Request

```
GET /orders/9876 HTTP/1.1  
Host: restbucks.com  
If-None-Match: "4d3e88c9"
```

Response

```
HTTP/1.1 200 OK  
ETag: "5612cfaa"  
Content-Type: application/restbucks+xml
```

```
<shop xmlns="http://schemas.restbucks.com/shop"  
      xmlns:rb="http://relations.restbucks.com/">  
    <status>Paid</status>  
</shop>
```

Server-driven content negotiation

Server chooses best representation based on information in the request

Request

GET /updates HTTP/1.1

Host: restbucks.com

Accept: application/atom+xml, application/xhtml+xml

Server-driven content negotiation

Server chooses best representation based on information in the request

Response

HTTP/1.1 200 OK

Cache-Control: max-age=10

Date: Fri, 29 Oct 2010 14:40:45 GMT

Server: Microsoft-IIS/6.0

Vary: Accept

ETag: "081fcddc"

Content-Location: <http://restbucks.com/updates.atom>

Content-Type: application/atom+xml

Content-Length: 5348

Last-Modified: Fri, 29 Oct 2010 14:40:44 GMT

```
<?xml version="1.0" encoding="utf-8"?>
<feed xmlns="http://www.w3.org/2005/Atom">
    ...
</feed>
```

Server-driven content negotiation (continued)

Selection criteria

Accept, Accept-Charset, Accept-Encoding, Accept-Language
Server can use other information in the request (e.g. User-Agent)

Vary header

Indicates which parameters were used to select the representation
Caches can use Vary to match cached responses with future requests
Beware of varying on headers with wildly differing values (e.g. Cookie)

Accept q parameter

Indicates client weightings
Accept: application/xml; q=0.5, application/atom+xml

Identifying variants

Sometimes useful to identify variants with Content-Location header

Agent-driven content negotiation

Client chooses best representation based on server-supplied list

Request

```
GET /documentation HTTP/1.1  
Host: restbucks.com
```

Response

```
HTTP/1.1 300 Multiple Choices  
Cache-Control: public  
ETag: "c549995b"  
Content-Length: 0  
Link: </documentation.txt>;type=text/plain,  
      </documentation.html>;type=text/html,  
      </documentation.pdf>;type=application/pdf
```

Request

```
GET /documentation.html HTTP/1.1  
Host: restbucks.com
```

Response

```
HTTP/1.1 200 OK  
...
```

Reliability by convention

Safe

No side-effects for which client is responsible

Idempotent

OK to repeat in event of failures

Verb	Safe?	Idempotent?	Description
GET	■	■	Retrieve representation
PUT		■	Store representation
POST			Create resource, handle data, annotate resource
DELETE		■	Delete resource

Operation-Oriented vs. Resource-Oriented Design

CreateQuote

SubmitQuote

AmendOrder

AddItem

SearchOrders

ReserveItem

CancelOrder

RemoveItem

Operations

Resources

GET
PUT
POST
DELETE

Quote

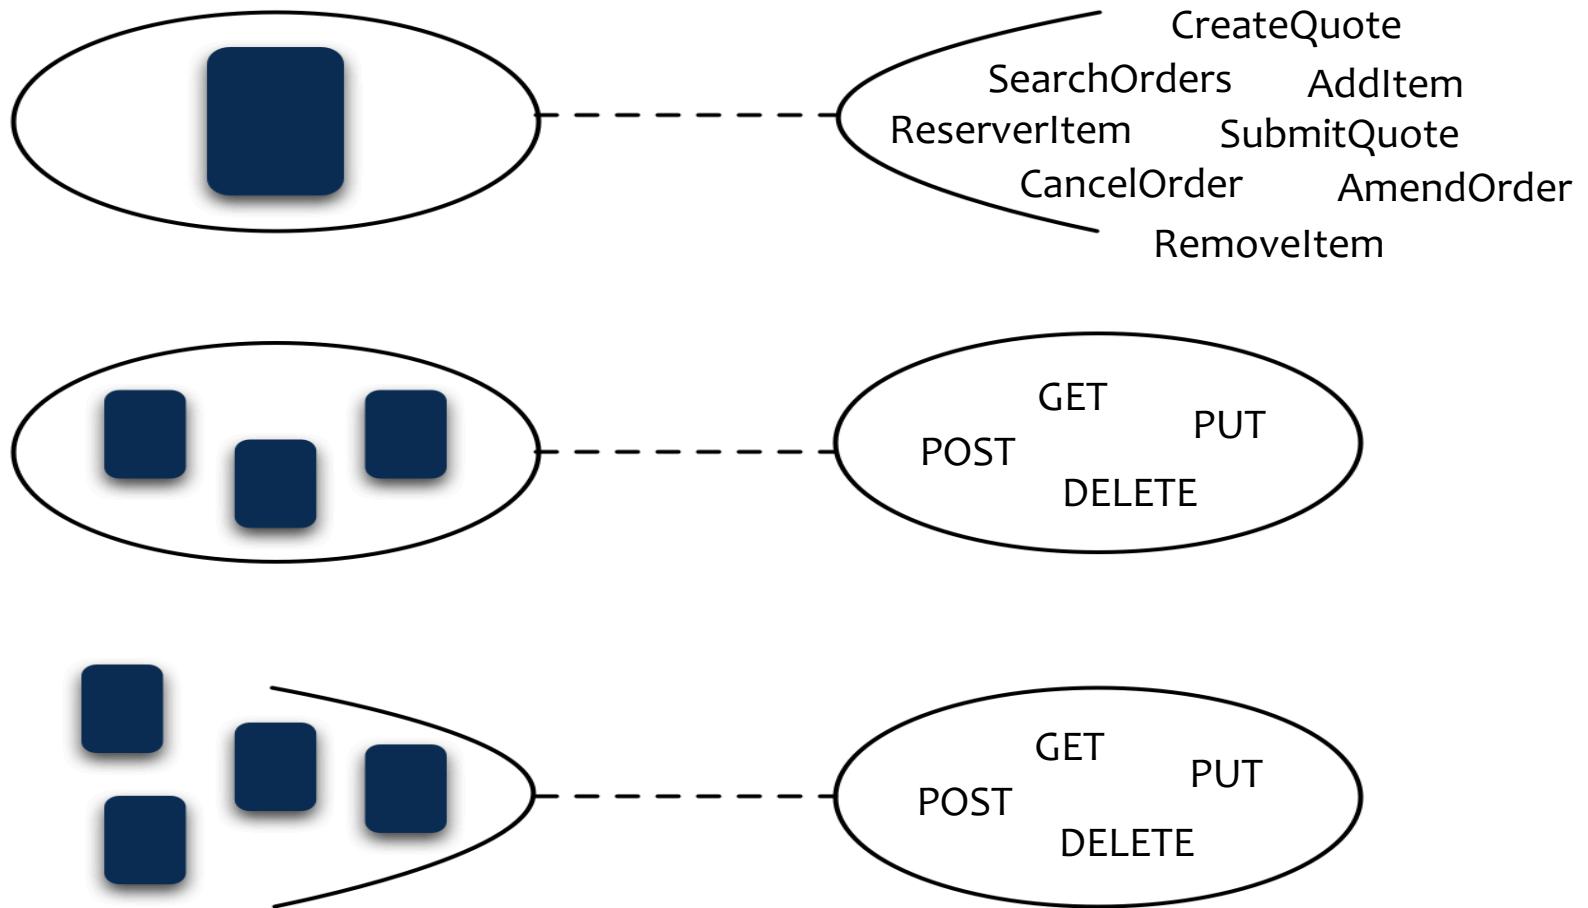
GET
PUT
POST
DELETE

Order

GET
PUT
POST
DELETE

Item

Specialization and Innovation Depend on an Open Set



HTTP/2

- Binary rather than textual protocol
 - compact, less error prone than textual
- Fully multiplexed, instead of ordered and blocking
 - Eliminates "head-of-line blocking" with large responses
- One connection for parallelism
 - Mitigate against buffers overflows in network, congestion and retransmits
- Header compression
 - Reduce number of TCP round trips to get headers onto the wire
- Server push
 - Allows servers to “push” responses proactively into client caches

Summary

- Web is a stateless, client-server, layered architecture
 - Global reach and evolvability achieved through uniform, coarse-grained interactions (rather than fine-grained, application-specific operations)
- Clients manipulate resources by exchanging documents (resource representations)
 - Resources identified and addressed using URIs
 - Resource representations use standardised media types
- Document coordination achieved using a document transfer application protocol (HTTP)

Tech Interlude

URI Templates



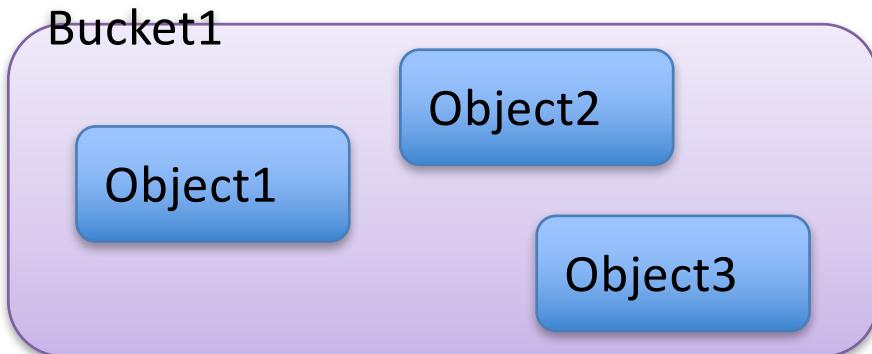
Conflicting URI Philosophies

- URIs should be descriptive, predictable?
 - `http://spreadsheet/cells/a2,a9`
 - `http://jim.webber.name/2007/06.aspx`
 - Convey some ideas about how the underlying resources are arranged
 - Can infer `http://spreadsheet/cells/b0,b10` and `http://jim.webber.name/2005/05.aspx` for example
 - Nice for programmatic access, but may introduce coupling
- URIs should be opaque?
 - `http://tinyurl.com/6`
 - TimBL says “opaque URIs are cool”
 - Convey no semantics, can’t infer anything from them
 - Don’t introduce coupling

URI Templates, in brief

- Use URI templates to make your resource structure easy to understand
- For Amazon S3 (storage service) it's easy:

`http://s3.amazonaws.com/{bucket-name}/{object-name}`



URI Templates are Easy!

- Take the URI:

`http://restbucks.com/orders?{order_id}`

- You could do the substitution and get a URI:

`http://restbucks.com/orders?1234`

- Can easily make more complex URIs too

- Mixing template and non-template sections

`http://restbucks.com/{orders}/{shop}/{year}/{month}.xml`

- Use URI templates client-side to compute server-side URIs
 - But beware this introduces coupling!

Why URI Templates?

- Regular URIs are a good idiom in Web-based services
 - Helps with understanding, self documentation
- They allow users to infer a URI
 - If the pattern is regular
- URI templates formalise this arrangement
 - And advertise a template rather than a regular URI

URI Templates Pros and Cons

- Everything interesting in a Web-based service has a URI
- Remember, two schools of thought:
 - Opaque URIs are cool (Berners-Lee)
 - Transparent URIs are cool (everyone else)
- URI templates present two core concerns:
 - They invite clients to invent URIs which may not be honoured by the server
 - They increase coupling since servers must honour forever any URI templates they've advertised
- Use URI templates sparingly, and with caution
 - Entry point URIs only is a good rule of thumb

Embracing HTTP: CRUD



Domain Overview: Restbucks

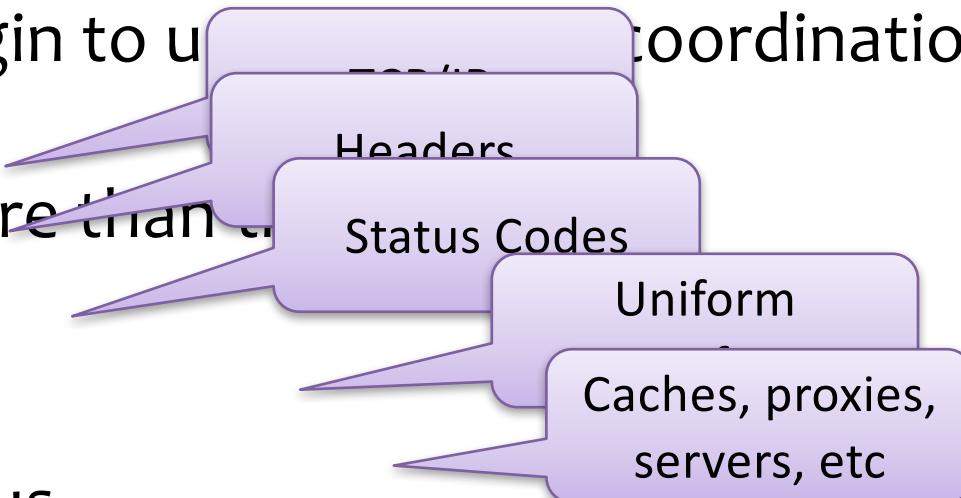
- Throughout *REST in Practice* we used an example drawn from conventional commerce
- A coffee shop called RESTbucks*
 - Influenced by Gregor Hohpe’s work “Startbucks does not use Two Phase Commit”
- Everybody knows how a coffee shop works, so the domain doesn’t get in the way
- Although the example is a little contrived, it has plenty of scope for dealing with real-world systems-of-microservices
- Let’s start with the protocol for order management...

CRUD APIs

- This session is about CRUD/HTTP fundamentals
- We see how HTTP is used to manipulate an object in a remote service
- But we also consider the limitations:
 - Coupling (design time, runtime)
 - Limited protocol expressivity
 - We'll explore these with a hands-on exercise

Using the Web

- SOAP services use the Web as a transport
- CRUD services begin to use the Web's coordination support
- But the Web is more than a transport
 - Transport, plus
 - Metadata, plus
 - Fault model, plus
 - Component model, plus
 - Runtime environment, plus...

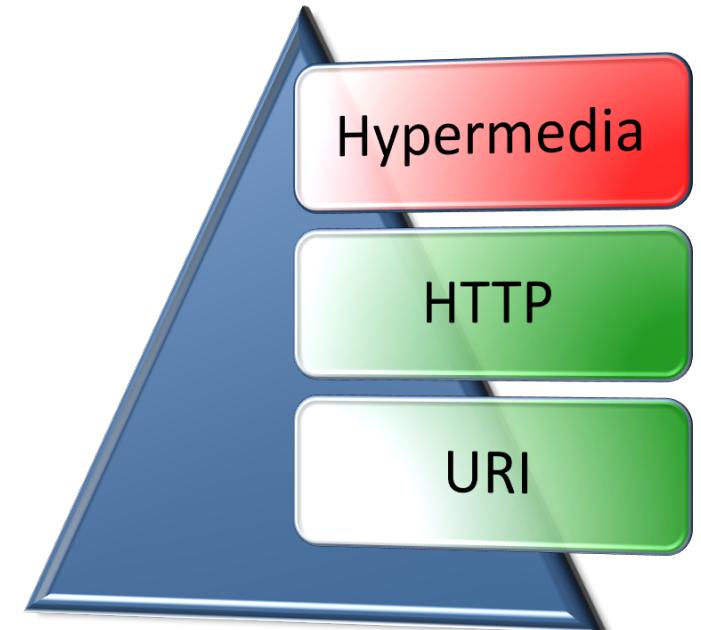


CRUD Resource Lifecycle

- The resource is created with POST
- It's read with GET
- And updated via PUT
- Finally it's removed using DELETE

Richardson Model Level 2

- Lots of URIs
- Understands HTTP!
- No hypermedia



Create with POST



POST Semantics

- POST creates a new resource
- But the server decides on that resource's URI
- Common human Web example: posting to Web log
 - Server decides URI of posting and any comments made on that post
- Programmatic Web example: creating a new employee record
 - And subsequently adding to it

POST Request

```
POST /orders HTTP/1.1
```

Verb, path, and HTTP
version

```
Host: restbucks.example.com
```

```
Content-Type:application/vnd.restbucks+xml
```

```
Content-Length: 225
```

Restbucks-specific XML
content

```
<order xmlns="http://schemas.restbucks.com">
```

Content

```
    <location>takeAway</location>
```

(again Restbucks XML)

```
    <item>
```

```
        <drink>latte</drink>
```

```
        <milk>whole</milk>
```

```
        <size>small</size>
```

```
    </item>
```

```
</order>
```

POST Response

HTTP/1.1 201 Created

Location: /orders/1234

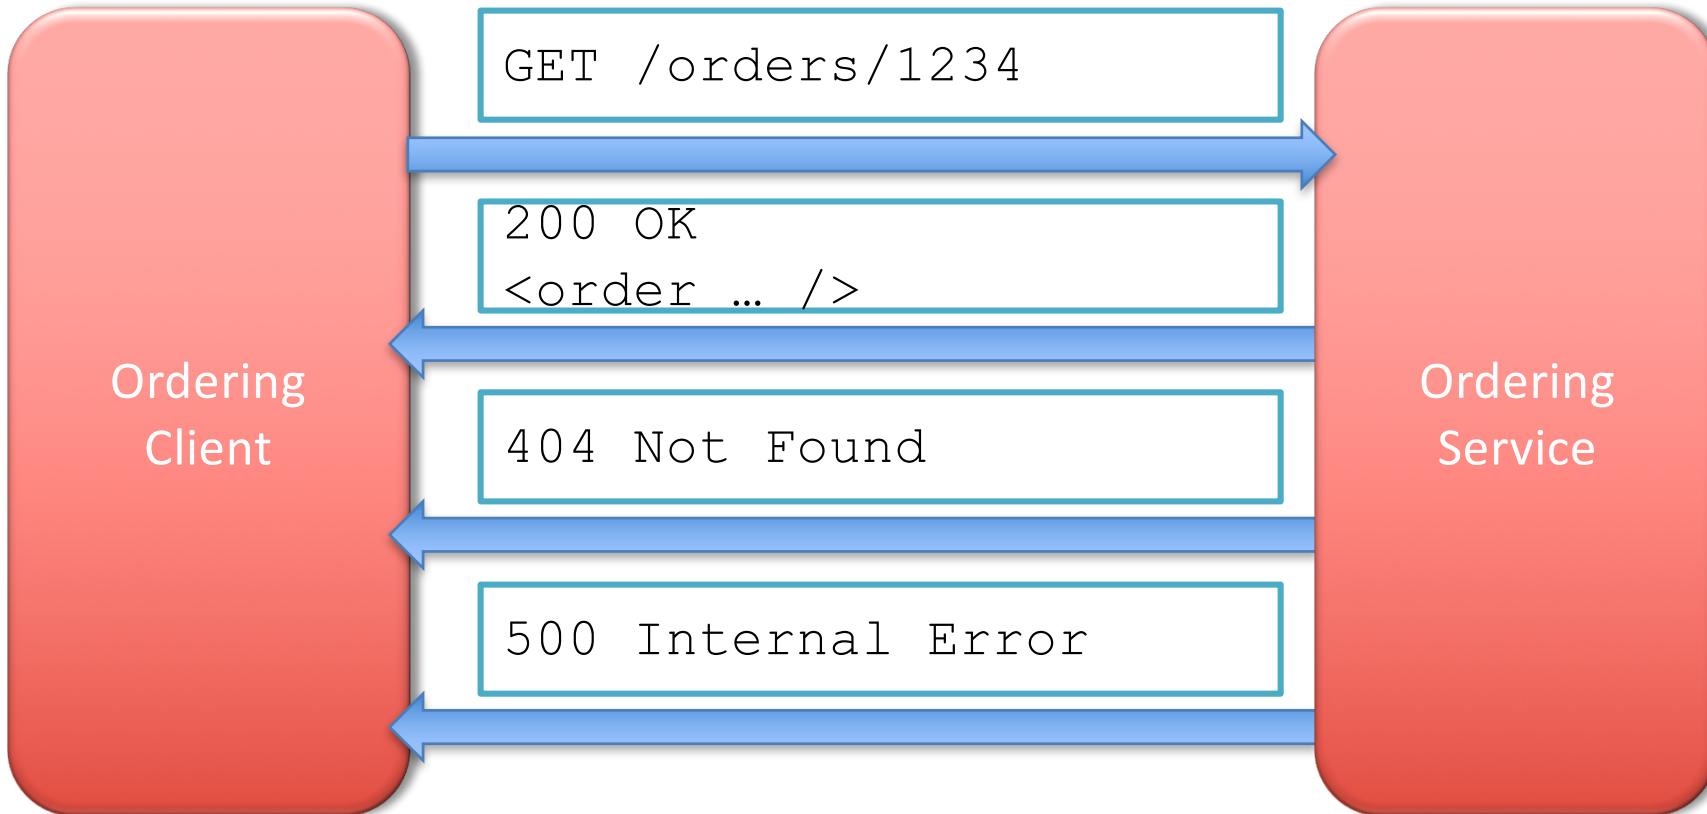
When POST goes wrong

- We may be 4xx or 5xx errors
 - Client versus server problem
- We turn to GET!
- Find out the resource states first
 - Then figure out how to make forward or backward progress
- Then solve the problem
 - May involve POSTing again
 - May involve a PUT to rectify server-side resources in-place
 - PUT coming soon

POST Implementation with a Servlet

```
protected void doPost(HttpServletRequest request,
                      HttpServletResponse response) {
    try {
        Order order = extractOrderFromRequest(request);
        String internalOrderId = OrderDatabase.getDatabase().saveOrder(order);
        response.setHeader("Location", computeLocationHeader(request,
                                                               internalOrderId));
        response.setStatus(HttpServletResponse.SC_CREATED);
    } catch (Exception ex) {
        response.setStatus(HttpServletResponse.SC_INTERNAL_SERVER_ERROR);
    }
}
```

Read with GET



GET Semantics

- GET retrieves the representation of a resource
- Should be idempotent
 - Shared understanding of GET semantics
 - Don't violate that understanding!

Library of congress
catalogue incident!

GET Exemplified

GET /orders/1234 HTTP/1.1

Accept: application/vnd.restbucks+xml

Host: restbucks.com

GET Response

HTTP/1.1 200 OK

Content-Length: 232

Content-Type: application/vnd.restbucks+xml

Date: Wed, 19 Nov 2008 21:48:10 GMT

```
<order xmlns="http://schemas.restbucks.com/order">
  <location>takeAway</location>
  <items>
    <item>
      <name>latte</name>
      <quantity>1</quantity>
      <milk>whole</milk>
      <size>small</size>
    </item>
  </items>
  <status>pending</status>
</order>
```

When GET Goes wrong

- Simple!
 - Just 404 – the resource is no longer available

HTTP/1.1 404 Not Found

Content-Type: application/vnd.restbucks+xml

Content-Length: 952

Date: Sat, 20 Dec 2008 19:01:33 GMT

- Are you sure?
 - GET again!
- GET is safe and idempotent
 - Great for crash recovery scenarios!

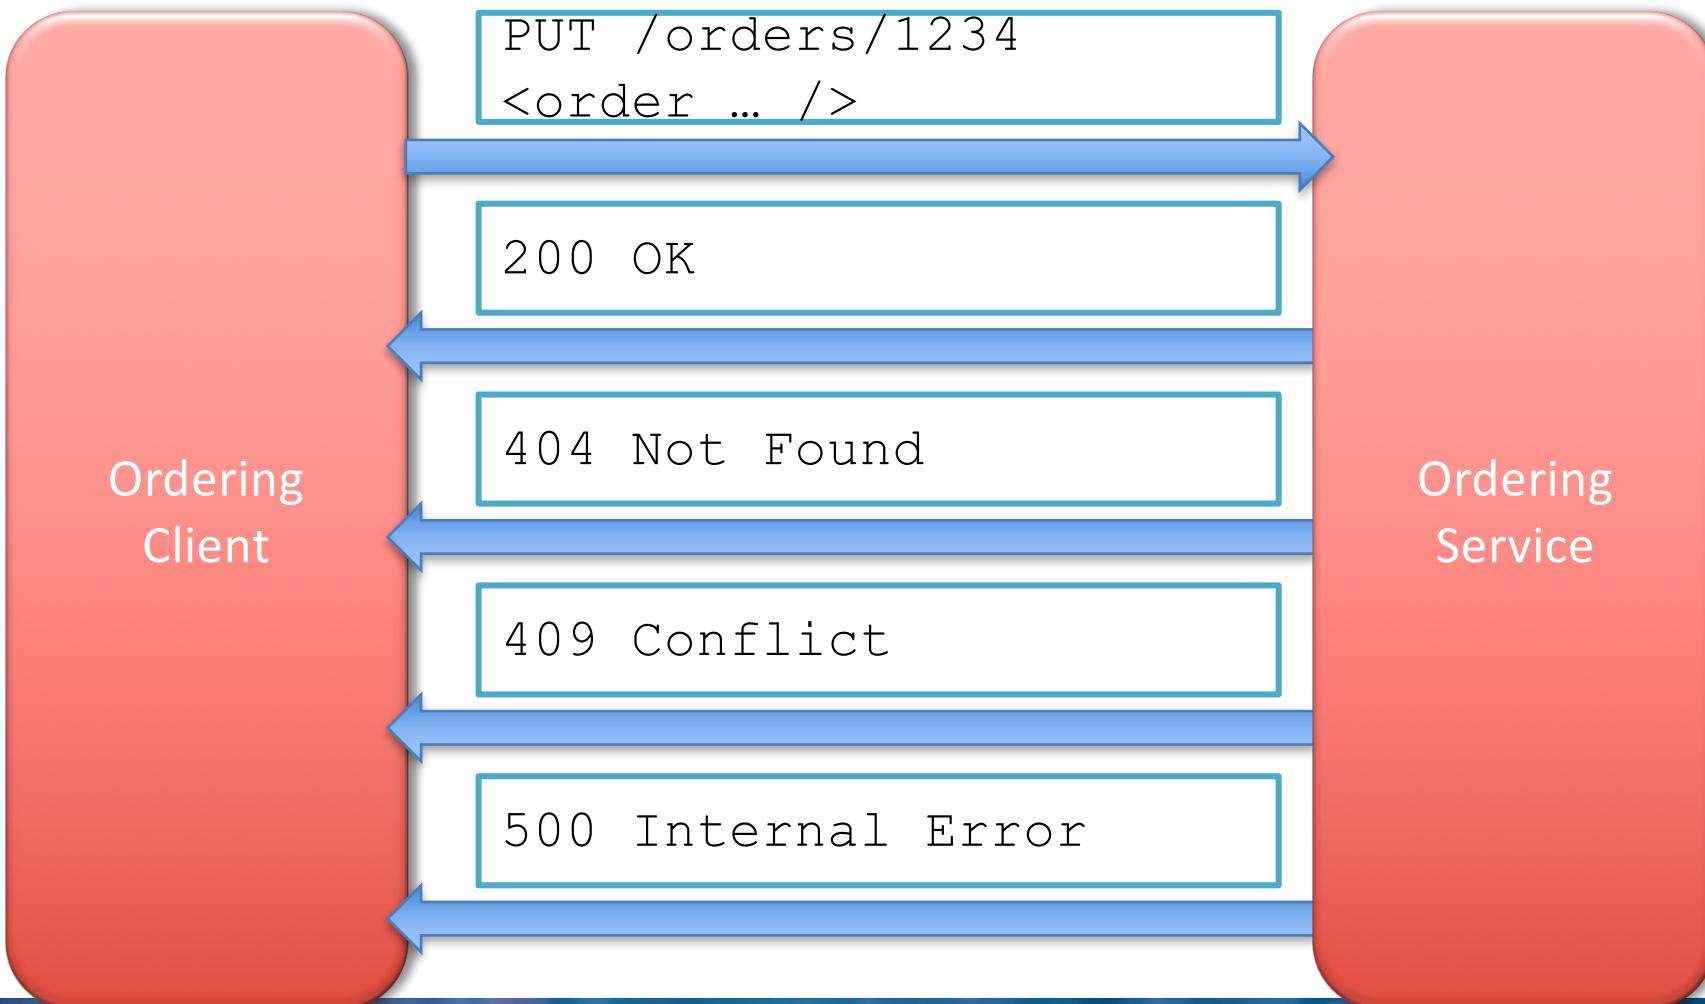
Idempotent Behaviour

- An action with no side affects
 - Comes from mathematics
- In practice means two things:
 - A safe operation is one which changes no state at all
 - E.g. HTTP GET
 - An idempotent operation is one which updates state in an absolute way
 - E.g. $x = 4$ rather than $x += 2$
- Web-friendly systems scale because of safety
 - Caching!
- And are fault tolerant because of idempotent behaviour
 - Just re-try in failure cases

GET JAX-RS Implementation

```
@Path("/")
public class OrderingService {
    @GET
    @Produces("application/xml")
    @Path("/{orderId}")
    public String getOrder(@PathParam("orderId") String orderId) {
        try {
            Order order = OrderDatabase.getDatabase().getOrder(orderId);
            if (order != null) {
                return xstream.toXML(order);
            } else {
                throw new WebApplicationException(404);
            }
        } catch (Exception e) {
            throw new WebApplicationException(500);
        }
    }
    // Remainder of implementation omitted for brevity
}
```

Update with PUT



PUT Semantics

- PUT creates a new resource but the client decides on the URI
 - Providing the server logic allows it
- Also used to update existing resources by overwriting them in-place
- PUT is idempotent
 - Makes absolute changes
- But is not safe
 - It changes state!

PUT Request

```
PUT /orders/1234 HTTP/1.1
```

```
Host: restbucks.com
```

```
Content-Type: application/xml
```

```
Content-Length: 386
```

```
<order xmlns="http://schemas.restbucks.com/order">
  <location>takeAway</location>
  <items>
    <item>
      <milk>whole</milk>
      <name>latte</name>
      <quantity>2</quantity>
      <size>small</size>
    </item>
    <item>
      <milk>whole</milk>
      <name>cappuccino</name>
      <quantity>1</quantity>
      <size>large</size>
    </item>
  </items>
  <status>preparing</status>
</order>
```

Updated content

PUT Response

HTTP/1.1 200 OK

Date: Sun, 30 Nov 2008 21:47:34 GMT

Content-Length: 0

Minimalist response contains no entity body

When PUT goes wrong

- If we get 5xx error, or some 4xx errors simply PUT again!
 - PUT is idempotent
- If we get errors indicating incompatible states (409, 417) then do some forward/backward compensating work
 - And maybe PUT again

```
HTTP/1.1 409 Conflict
Date: Sun, 21 Dec 2008 16:43:07 GMT
Content-Length:382

<order xmlns="http://schemas.restbucks.com/order">
  <location>takeAway</location>
  <items>
    <item>
      <milk>whole</milk>
      <name>latte</name>
      <quantity>2</quantity>
      <size>small</size>
    </item>
    <item>
      <milk>whole</milk>
      <name>cappuccino</name>
      <quantity>1</quantity>
      <size>large</size>
    </item>
  </items>
  <status>served</status>
</order>
```

WCF Implementation for PUT

```
[ServiceContract]  
public interface IOrderingService  
{  
    [OperationContract]  
    [WebInvoke(Method = "PUT", UriTemplate =  
        "/orders/{orderId}")]  
    void UpdateOrder(string orderId, Order order);  
  
    // ...  
}
```

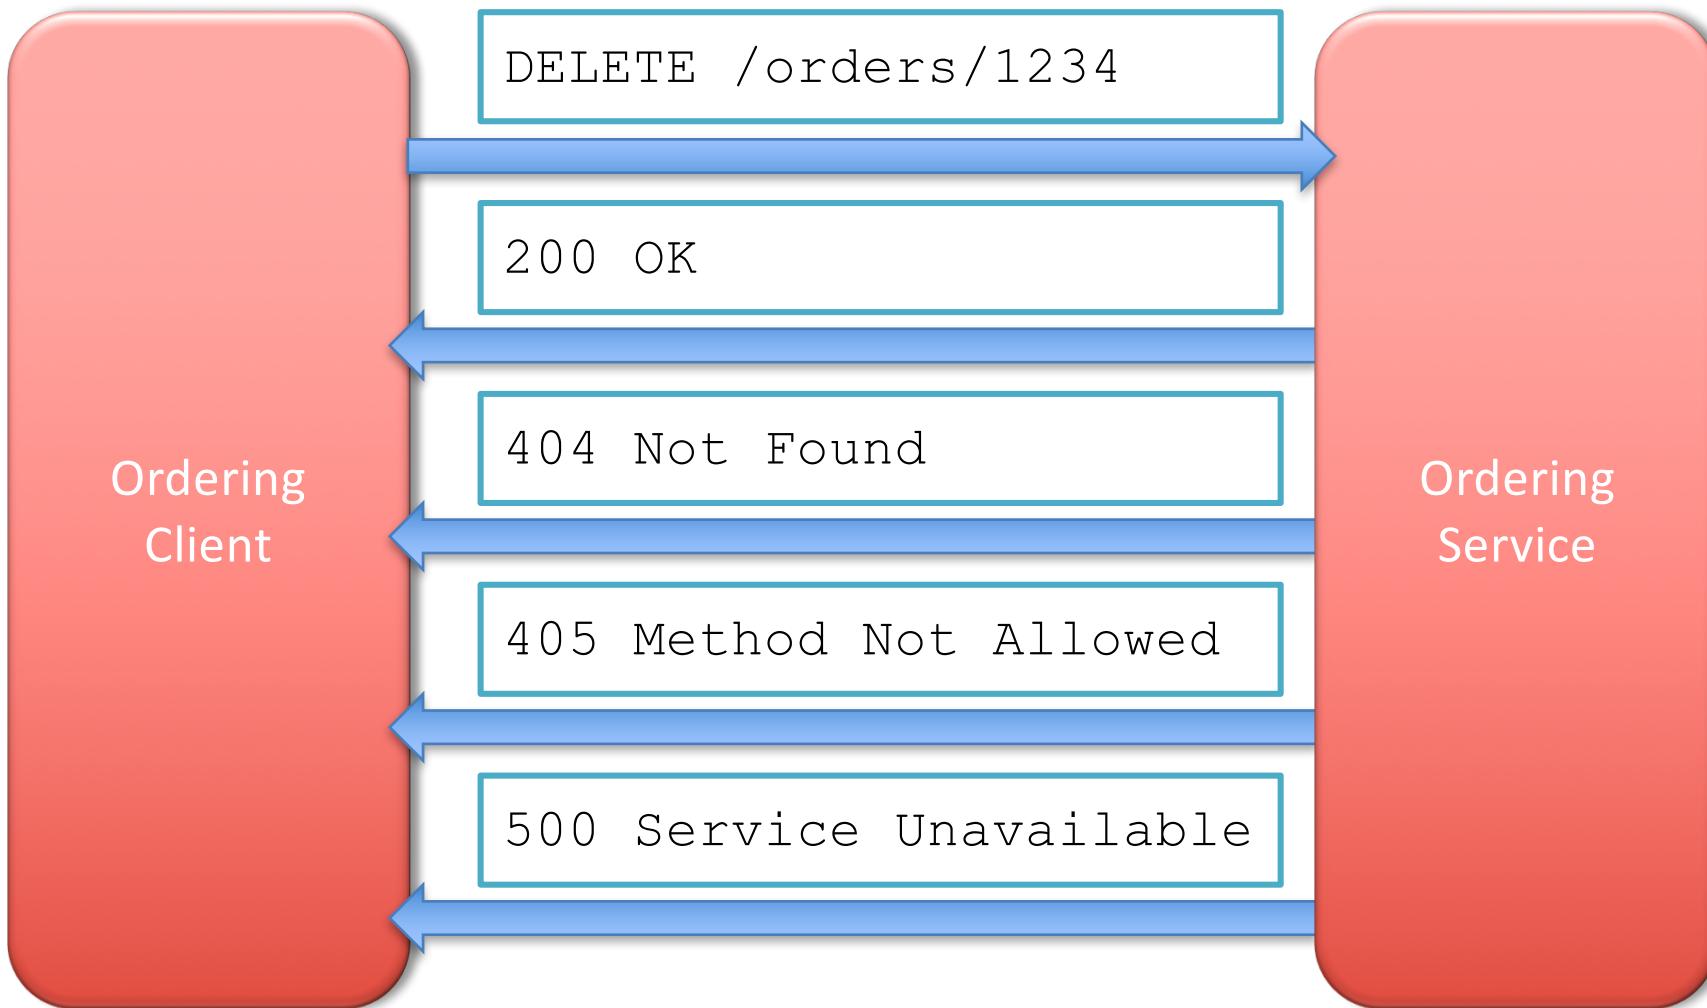
WCF Serializable Types

```
[DataContract(Namespace = "http://schemas.restbucks.com/order", Name = "order")]
public class Order
{
    [DataMember(Name = "location")]
    public Location ConsumeLocation
    {
        get { return location; }
        set { location = value; }
    }

    [DataMember(Name = "items")]
    public List<Item> Items
    {
        get { return items; }
        set { items = value; }
    }

    [DataMember(Name = "status")]
    public Status OrderStatus
    {
        get { return status; }
        set { status = value; }
    }
    // ...
}
```

Remove with DELETE



DELETE Semantics

- Stop the resource from being accessible
 - Logical delete, not necessarily physical
- Request

```
DELETE /orders/1234 HTTP/1.1
```

```
Host: restbucks.com
```

- Response

```
HTTP/1.1 200 OK
```

```
Content-Length: 0
```

```
Date: Tue, 16 Dec 2008 17:40:11 GMT
```

This is important for
decoupling
implementation details
from resources

When DELETE goes wrong

- Simple case, DELETE again!
 - DELETE is idempotent!
 - DELETE once, DELETE 10 times has the same effect:
one deletion

HTTP/1.1 404 Not Found

Content-Length: 0

Date: Tue, 16 Dec 2008 17:42:12 GMT

When DELETE goes Really Wrong

- Some 4xx responses indicate that deletion isn't possible
 - The state of the resource isn't compatible
 - Try forward/backward compensation instead

```
HTTP/1.1 409 Conflict
Content-Length: 379
Date: Tue, 16 Dec 2008 17:53:09 GMT

<order xmlns="http://schemas.restbucks.com/order">
  <location>takeAway</location>
  <items>
    <item>
      <name>latte</name>
      <milk>whole</milk>
      <size>small</size>
      <quantity>2</quantity>
    </item>
    <item>
      <name>cappuccino</name>
      <milk>skim</milk>
      <size>large</size>
      <quantity>1</quantity>
    </item>
  </items>
  <status>served</status>
</order>
```

Can't delete an order that's already served



CRUD does not mean Worthless

Exercise

- Write a client for the service hosted on this laptop
- It's an ordering service and you can CRUD the state of an order
- Here are some URI templates you can use:
 - /
 - /order
 - /order/{order_id}
 - /receipt/{order_id}
 - /payment
 - /payment/{payment_id}

Network

- There's an ad-hoc network called RESTinPractice
- Join it
- Then see if you can point your browser at:
 - <http://169.254.147.226:8080>

Exercise

- Order XML* template you can use:

```
<order xmlns="http://schemas.restbucks.com">  
  <item>  
    <milk>{whole | semi | skim | none}</milk>  
    <size>{small | large | medium}</size>  
    <drink>  
      {espresso | latte | cappuccino | americano}  
    </drink>  
  </item>  
  <location>takeaway</location>  
</order>
```

application/vnd.restbucks+xml



Exercise

- Payment template

```
<payment xmlns="http://schemas.restbucks.com">  
    <amount>{x.y}</amount>  
    <cardholderName>{cardholder name}</cardholderName>  
    <cardNumber>{card number no spaces}</cardNumber>  
    <expiryMonth>{mm}</expiryMonth>  
    <expiryYear>{yy}</expiryYear>  
</payment>
```

Curl (if you don't want to “code”)

- Windows users:
 - <https://curl.haxx.se/download.html>
- Single-line example:

```
curl -X POST -d @order.xml
http://localhost:8080/order
-H "Content-Type:
application/vnd.restbucks+xml"
-H "Accept:
application/vnd.restbucks+xml" -v
```



THOUGHT WE WERE DOING JSON FOR MICROSERVICES

GOT XML

CRUD is Good?

- CRUD is good
 - But it's not great
- CRUD-style services use some HTTP features
- But the application model is limited
 - Suits database-style applications
 - Hence frameworks like Microsoft's Astoria
- CRUD has limitations
 - CRUD ignores hypermedia
 - CRUD encourages tight coupling through URI templates
 - CRUD encourages server and client to collude
- The Web supports more sophisticated patterns than CRUD!

Tech Interlude

Hypermedia Formats



Media Types Rule!

- The Web's contracts are expressed in terms of media types and link relations
 - If you know the type, you can process the content
- Some types are special because they work in harmony with the Web
 - We call these “hypermedia formats”

(Hyper) media types

Standardised media type

Processing model

Hypermedia controls
(links and forms)

Supported operations
(methods, headers and
status codes)

Representation formats
(may include schemas)

Compose application-specific behaviours on top of the handling of standardised media types

General



Specific

Other Resource Representations

- Remember, XML is not the only way a resource can be serialised
 - Remember the Web is based on Representational State Transfer
- The choice of representation is left to the implementer
 - Can be a standard registered media type
 - Or something else
- But there is a division on the Web between two families
 - Hypermedia formats
 - Formats which host URIs and links
 - Regular formats
 - Which don't

Plain Old XML is not Hypermedia Friendly

HTTP/1.1 200 OK

Content-Length: 227

Content-Type: application/xml

Date: Wed, 19 Nov 2008 21:48:10 GMT

```
<order xmlns="http://schemas.restbucks.com/order">
  <location>takeAway</location>
  <items>
    <item>
      <name>latte</name>
      <quantity>1</quantity>
      <milk>whole</milk>
      <size>small</size>
    </item>
  </items>
  <status>pending</status>
</order>
```

Where are the links?
Where's the protocol?

So what?

- How do you know the next thing to do?
- How do you know the resources you're meant to interact with next?
- In short, how do you know the service's protocol?
 - Turn to WADL? Yuck!
 - Read the documentation? Come on!
 - URI Templates? Tight Coupling!

URI Templates are NOT a Hypermedia Substitute

- Often URI templates are used to advertise all resources a service hosts
 - Do we really need to advertise them all?
- This is verbose
- This is out-of-band communication
- This encourages tight-coupling to resources through their URI template
- This has the opportunity to cause trouble!
 - Knowledge of “deep” URIs is baked into consuming programs
 - Services encapsulation is weak and consumers will program to it
 - Service will change its implementation and break consumers

Bad Ideas with URI Templates

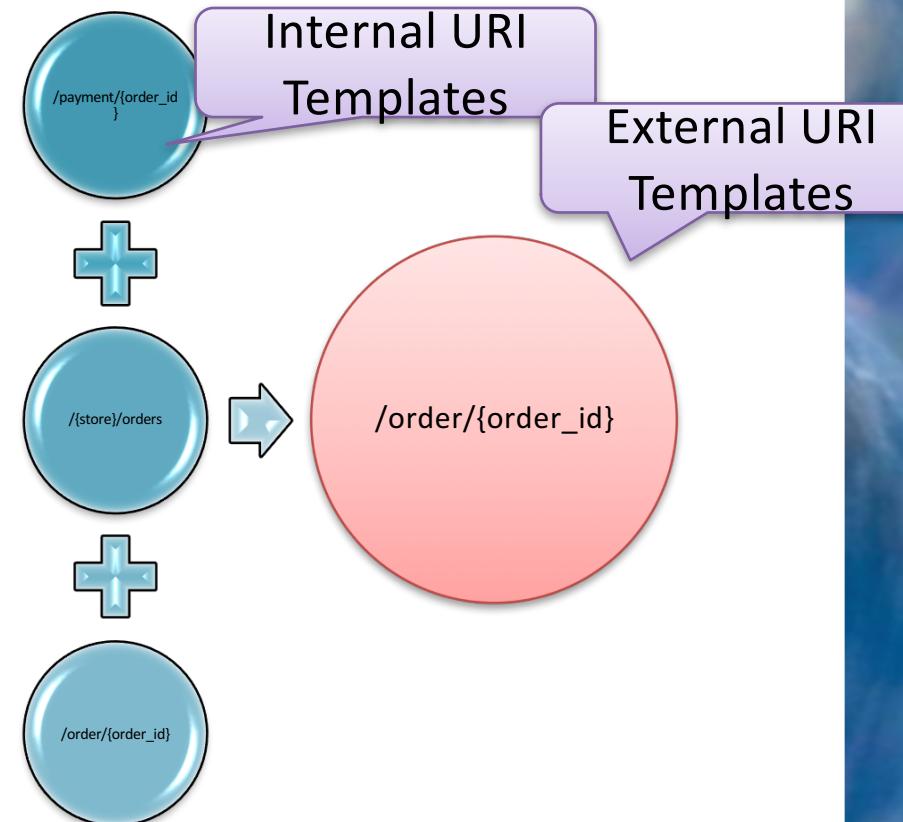
- Imagine we're created an order, what next?
- We could share this URI template:
 - `http://restbucks.com/payment/{order_id}`
- The `order_id` field should match the order ID that came from the `restbucks` service
 - Sounds great!
- But what if Restbucks outsources payment?
 - Change the URI for payments, break the template, break consumers!
 - D'oh!
- Be careful what you share!

Better Ideas for URI Templates: Entry Points

- Imagine that we have a well-known entry point to our service
 - Which corresponds to a starting point for a protocol
- Why not advertise that with a URI template?
- For example:
 - `http://restbucks.com/signIn/{store_id}/ {barista_id}`
- Changes infrequently
- Is important to Restbucks
- Is transparent, and easy to bind to

Better Ideas for URI Templates: Documentation!

- Services tend to support lots of resources
- We need a shorthand for talking about a large number of resources easily
- We can use a URI template for each “type” of resource that a service (or services) supports
- But we don’t share this information with others
 - Don’t violate encapsulation!



application/xml or application/json not enough

- Remember that HTTP is an application protocol
 - Headers and representations are intertwined
 - Headers set processing context for representations
- Remember that application/xml and application/json have specific processing models
 - Which doesn't include understanding the semantics of links
- Remember if a representation is declared in the Content-Type header, you must treat it that way
 - HTTP is an application protocol – did you forget already? ☺
- We need real hypermedia formats!

Hypermedia Formats

- Standard
 - Wide “reach”
 - Software agents already know how to process them
 - But sometimes need to be shoe-horned
- Self-created
 - Can craft specifically for domain
 - Semantically rich
 - But lack reach

Two Common Hypermedia Formats: XHTML and ATOM

- Both are commonplace today
- Both are hypermedia formats
 - They contain links
- Both have a processing model that explicitly supports links
- Which means both can describe protocols...

XHTML

- XHTML is just HTML that is also XML
- For example:

```
<html xmlns="http://www.w3.org/1999/xhtml"  
      xml:lang="en" lang="en"  
      xmlns:r="http://restbucks.org">  
  
  <head>  
    <title> XHTML Example </title>  
  </head>  
  
  <body>  
    <p>  
      . . .  
    </p>  
  </body>  
</html>
```

What's the big deal with XHTML?

- It does two interesting things:
 1. It gives us document structure
 2. It gives us links
- So?
 1. We can understand the format of those resources
 2. We can discover other resources!
- How?
 1. Follow the links!
 2. Encode the resource representation as “normal” XML in your XHTML documents
- Contrast this with the Atom and APP approach...similar!

XHTML in Action

```
<html xmlns="http://www.w3.org/1999/xhtml">
<body>
  <div class="order">
    <p class="location">takeAway</p>
    <ul class="items">
      <li class="item">
        <p class="name">latte</p>
        <p class="quantity">1</p>
        <p class="milk">whole</p>
        <p class="size">small</p>
      </li>
    </ul>
    <a href="http://restbucks.com/payment/1234"
       rel="payment">payment</a>
  </div>
</body>
</html>
```

Business data

“Hypermedia
Control”

application/xhtml+xml

- Can ask which verb the resource at the end of the link supports
 - Via HTTP OPTIONS
- No easy way to tell what each link actually does
 - Does it buy the music?
 - Does it give you lyrics?
 - Does it vote for a favourite album?
- We lack semantic understanding of the linkspace and resources
 - But we have microformats for that semantic stuff!
- Importantly XHTML is a hypermedia format
 - It contains hypermedia controls that can be used to describe protocols!

Atom Syndication Format

- We'll study this in more depth later, but for now...
- The application/atom+xml media type is hypermedia aware
- You should expect links when processing such representations
- And be prepared for a nascent protocol with them!

```
HTTP/1.1 200 OK
Content-Length: 342
Content-Type: application/atom+xml
Date: Sun, 22 Mar 2009 17:04:10 GMT
```

```
<entry xmlns="http://www.w3.org/2005/Atom">
  <title>Order 1234</title>
  <link rel="payment"
    href="http://restbucks.com/payment/1234"/>
  <link rel="special-offer"
    href="http://restbucks.com/offers/freeCookie"/>
  <id>http://restbucks.com/order/1234</id>
  <updated>2009-03-22T16:57:02Z</updated>
  <summary>1x Cafe Latte</summary>
  <try>
```

Links to other resources,
a nascent protocol

application/atom+xml

- No easy way to tell what each link actually does
 - But look at the way the `rel` attribute is being used
 - Can we inject semantics there?
- Atom is a hypermedia format
 - Both feeds and entries contains hypermedia controls that can describe protocols

Tech Interlude

Semantics



Microformats

- Microformats are an example of little “s” semantics
- Innovation at the edges of the Web
 - Not by some central design authority (e.g. W3C)
- Started by embedding machine-processable elements in Web pages
 - E.g. Calendar information, contact information, etc
 - Using existing HTML features like class, rel, etc

Semantic versus semantic

- Semantic Web is top-down
 - Driven by the W3C with extensive array of technology, standards, committees, etc
 - Has not currently proven as scalable as the visionaries hoped
 - RDF triples have been harvested and processed in private databases
- Microformats are bottom-up
 - Little formal organisation, no guarantee of interoperability
 - Popular formats tend to be adopted (e.g. hCard)
 - Easy to use and extend for our systems
 - Trivial to integrate into current and future programmatic Web systems

Microformats and Resources

- Use Microformats to structure resources where formats exist
 - I.e. Use hCard for contacts, hCalendar for data
- Create your own formats (sparingly) in other places
 - Annotating links is a good start
 - `<link rel="withdraw.cash" . . . />`
 - `<link rel="service.post" type="application/atom+xml" href="{post-uri}" title="some title">`
- The `rel` attribute describes the semantics of the referred resource

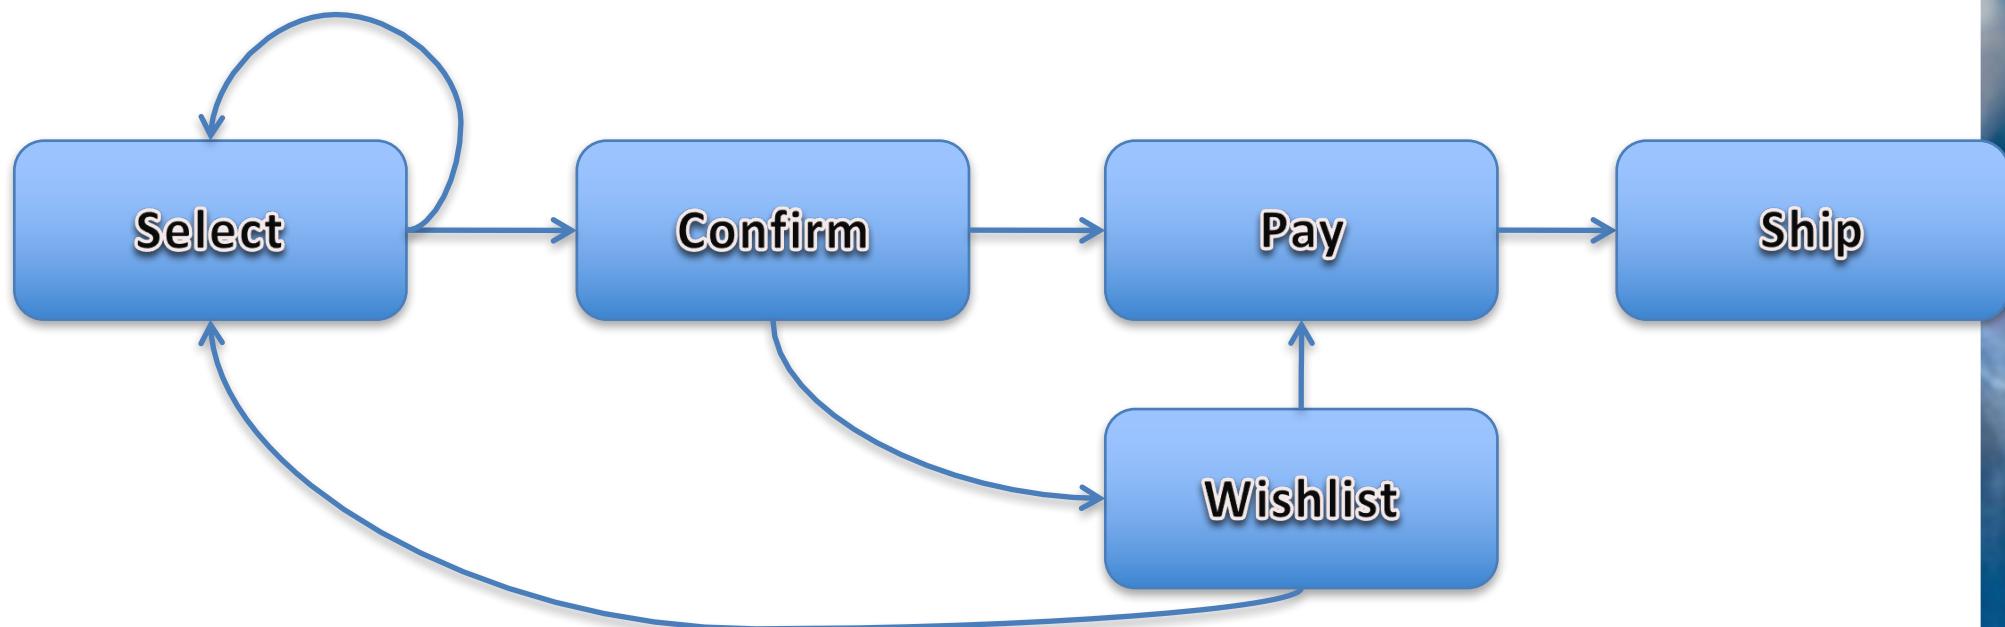
Hypermedia-Driven Microservices



Links

- Connectedness is good in Web-based systems
- Resource representations can contain other URIs
- Links advertise opportunities for state transitions
- Application (conversation) state is captured in terms of these states

Links advertise State Transitions



Links for APIs

```
<confirm xmlns="...">  
<link rel="payment"  
      href="https://pay"  
      type="application/xml"/>  
<link rel="postpone"  
      href="https://wishlist"  
      type="application/xml"/>  
</confirm>
```

- Following a link causes an action to occur
- This is the start of a state machine!
- Links lead to other resources which also have links
- Can make this stronger with semantics
 - Microformats

Anatomy of a link

```
<drink>
  <name>latte</name>
  <milk>whole</milk>
  <size>small</size>
  <link
    ○ href="http://restbucks.com/payment/1234"
    ○ rel="http://relations.restbucks.com/payment"
    ○ type="application/vnd.restbucks+xml"/>
</drink>
```

Address

WHERE does the linked resource reside?

Information processing

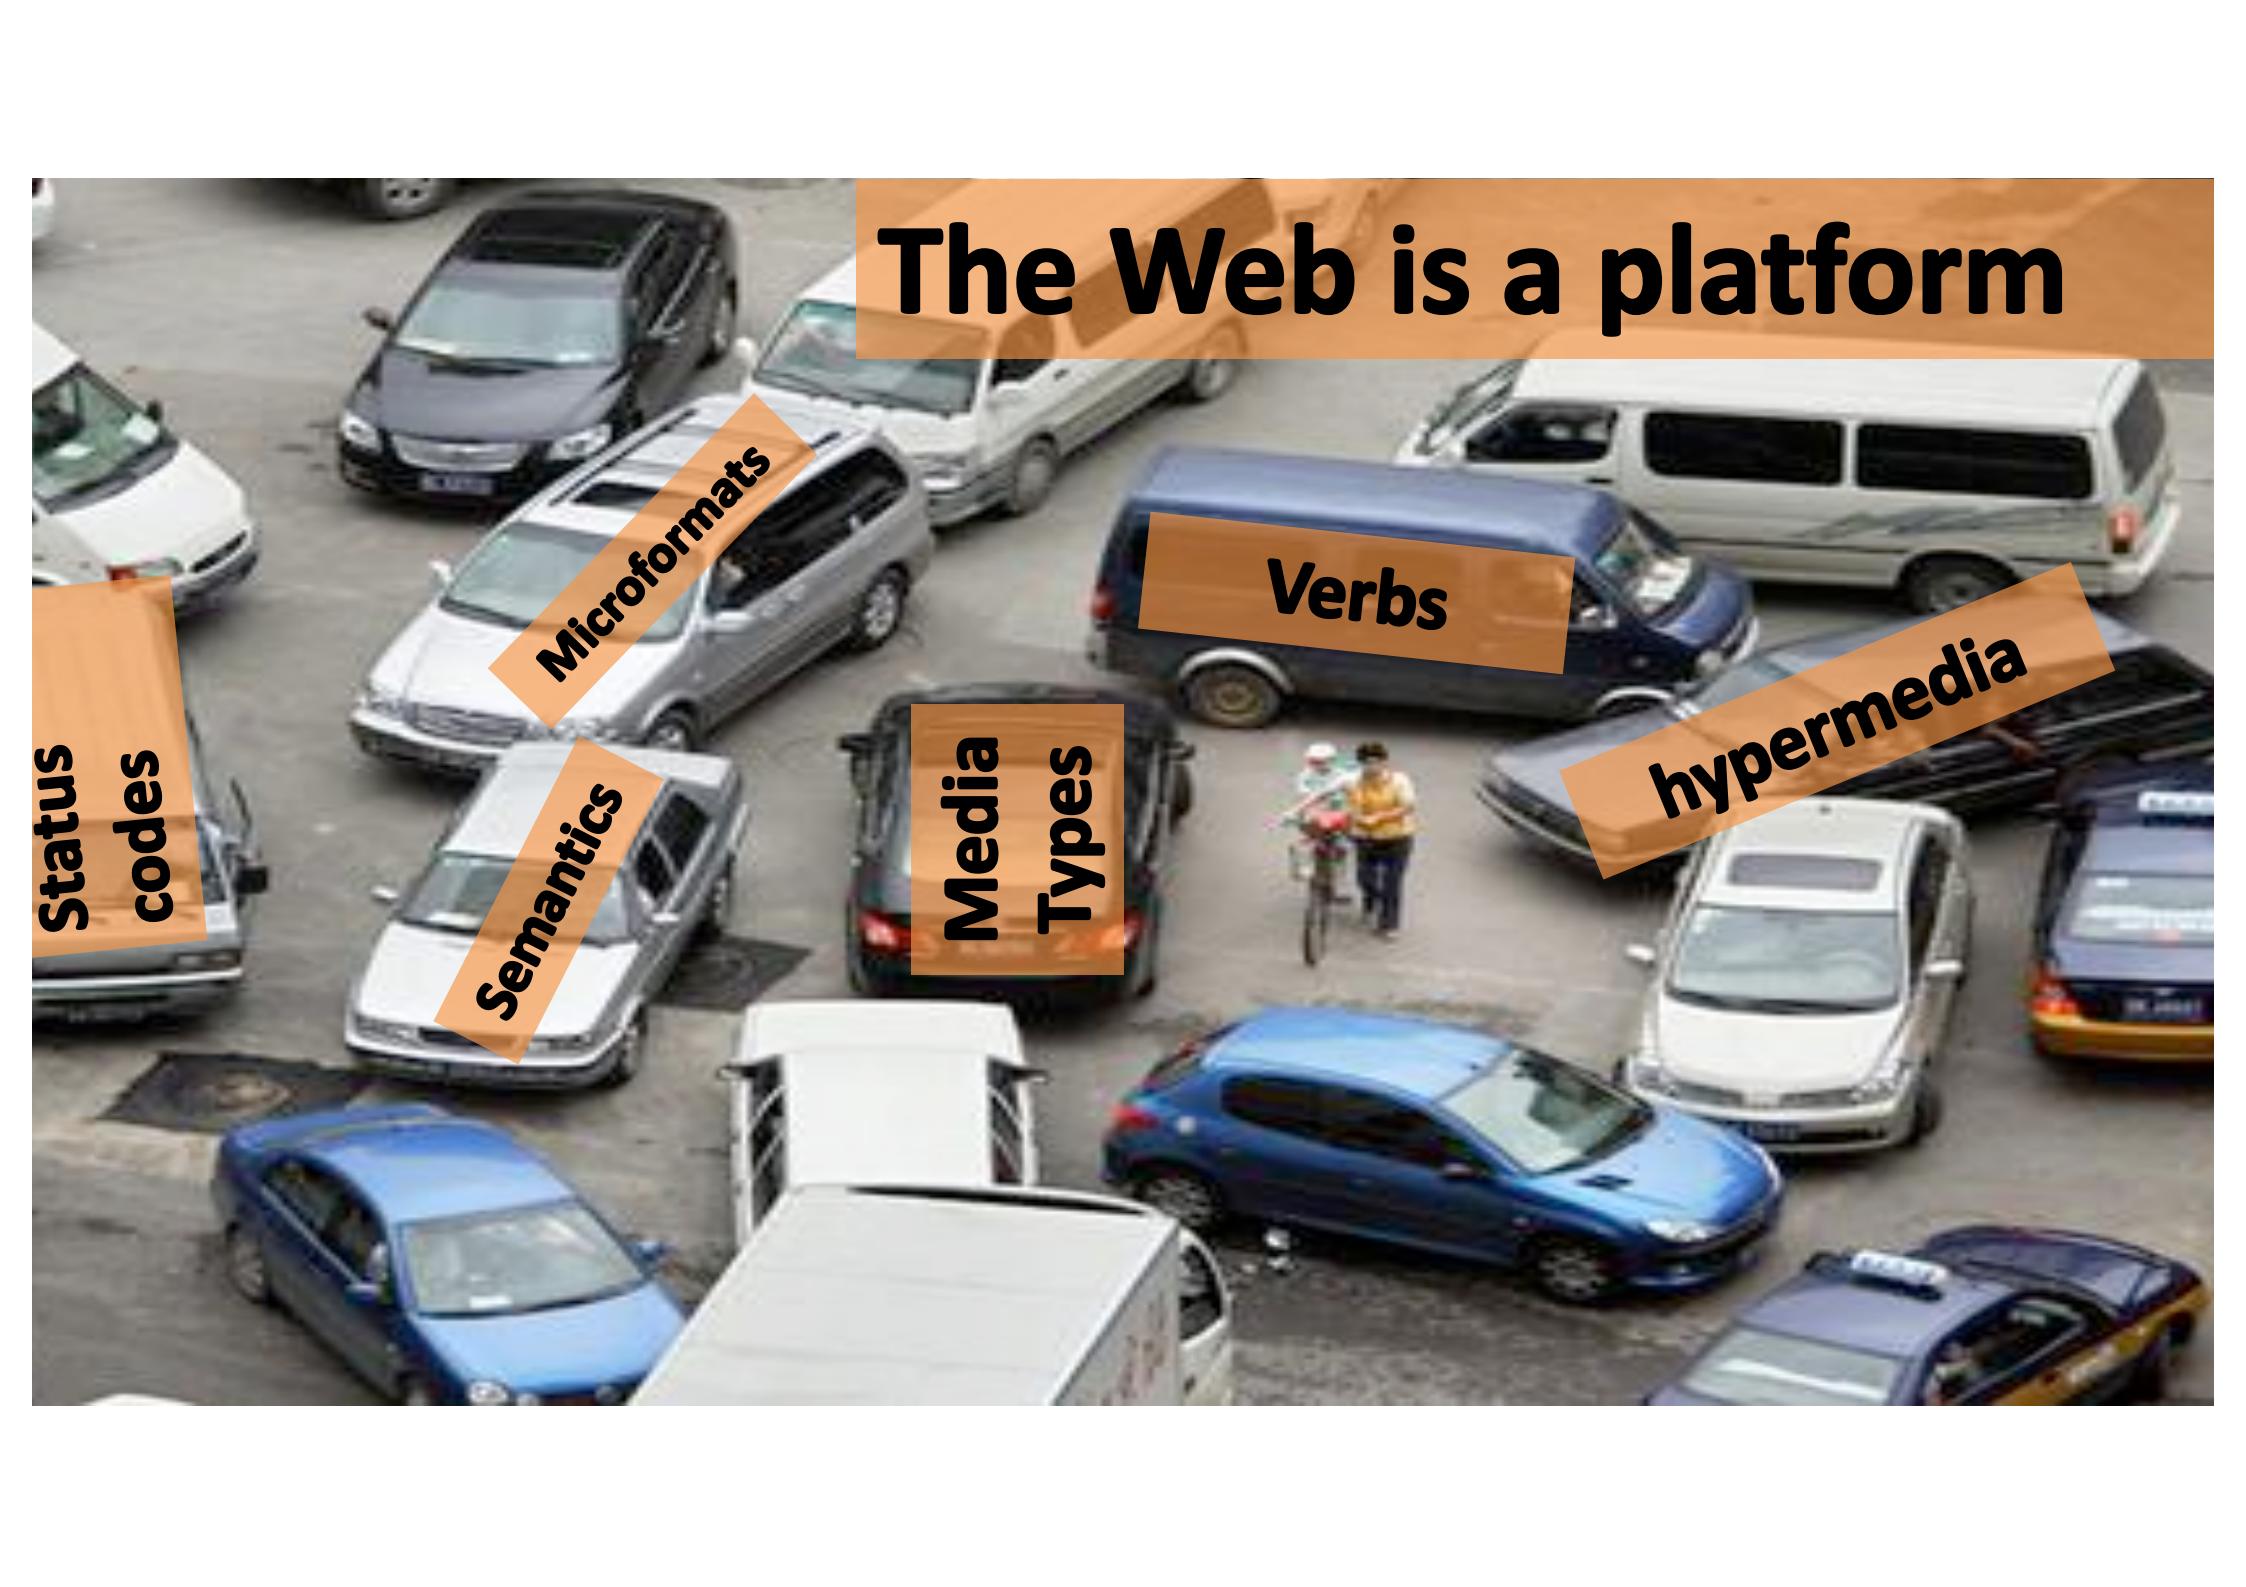
WHAT form does the linked resource take?

Semantic context

WHY access the linked resource?

We have a framework!

- The Web gives us a processing and metadata model
 - Verbs and status codes
 - Headers
- Gives us metadata contracts or Web “APIs”
 - URI Templates
 - Links
- Strengthened with semantics
 - Little “s”

An aerial photograph of a city street filled with various cars, including sedans, SUVs, and vans. A person is riding a bicycle in the center of the street. The image serves as the background for the slide.

The Web is a platform

Status
codes

Semantics

Microformats

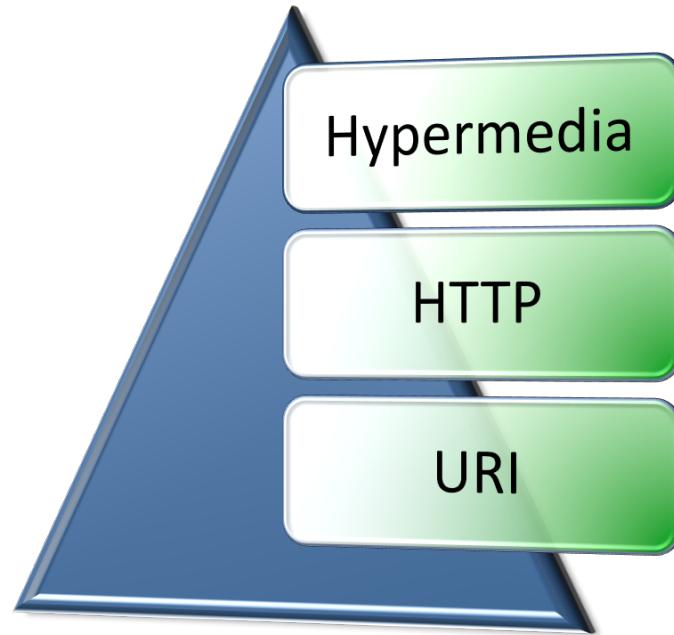
Media
Types

Verbs

hypermedia

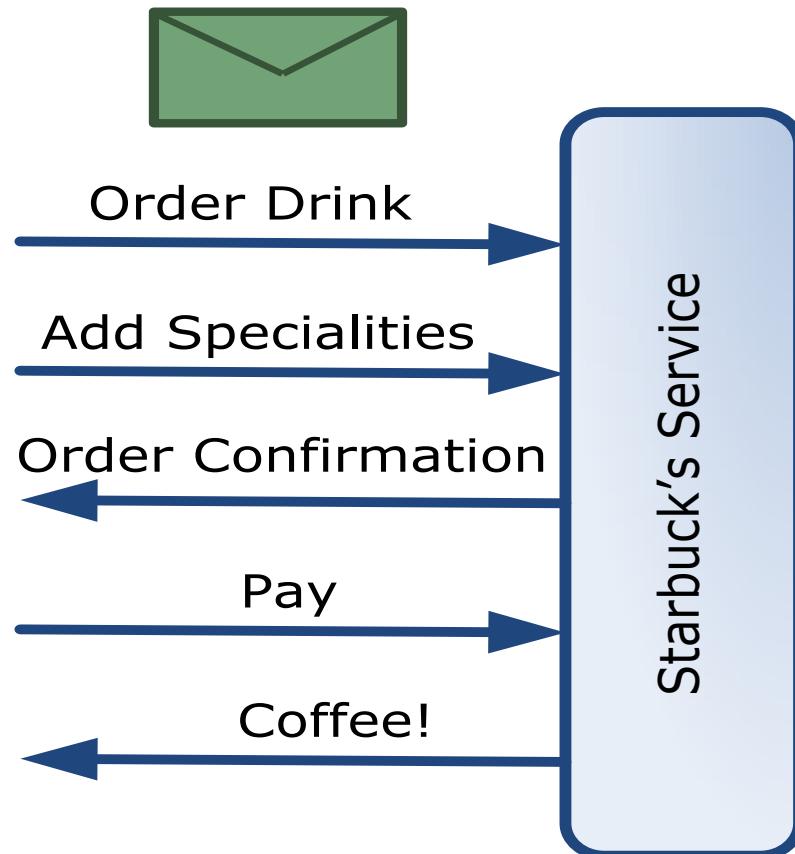
Richardson Model Level 3

- Lots of URIs that address resources
- Embraces HTTP as an application protocol
- Resource representations and formats identify other resources
 - Hypermedia at last!



Workflow and MOM

- With Web Services we exchange messages with the service
- Resource state is hidden from view
- Conversation state is all we know
 - Advertise it with SSDL, BPEL
- Uniform interface, roles defined by SOAP
 - No “operations”



Hypermedia Describes Protocols!

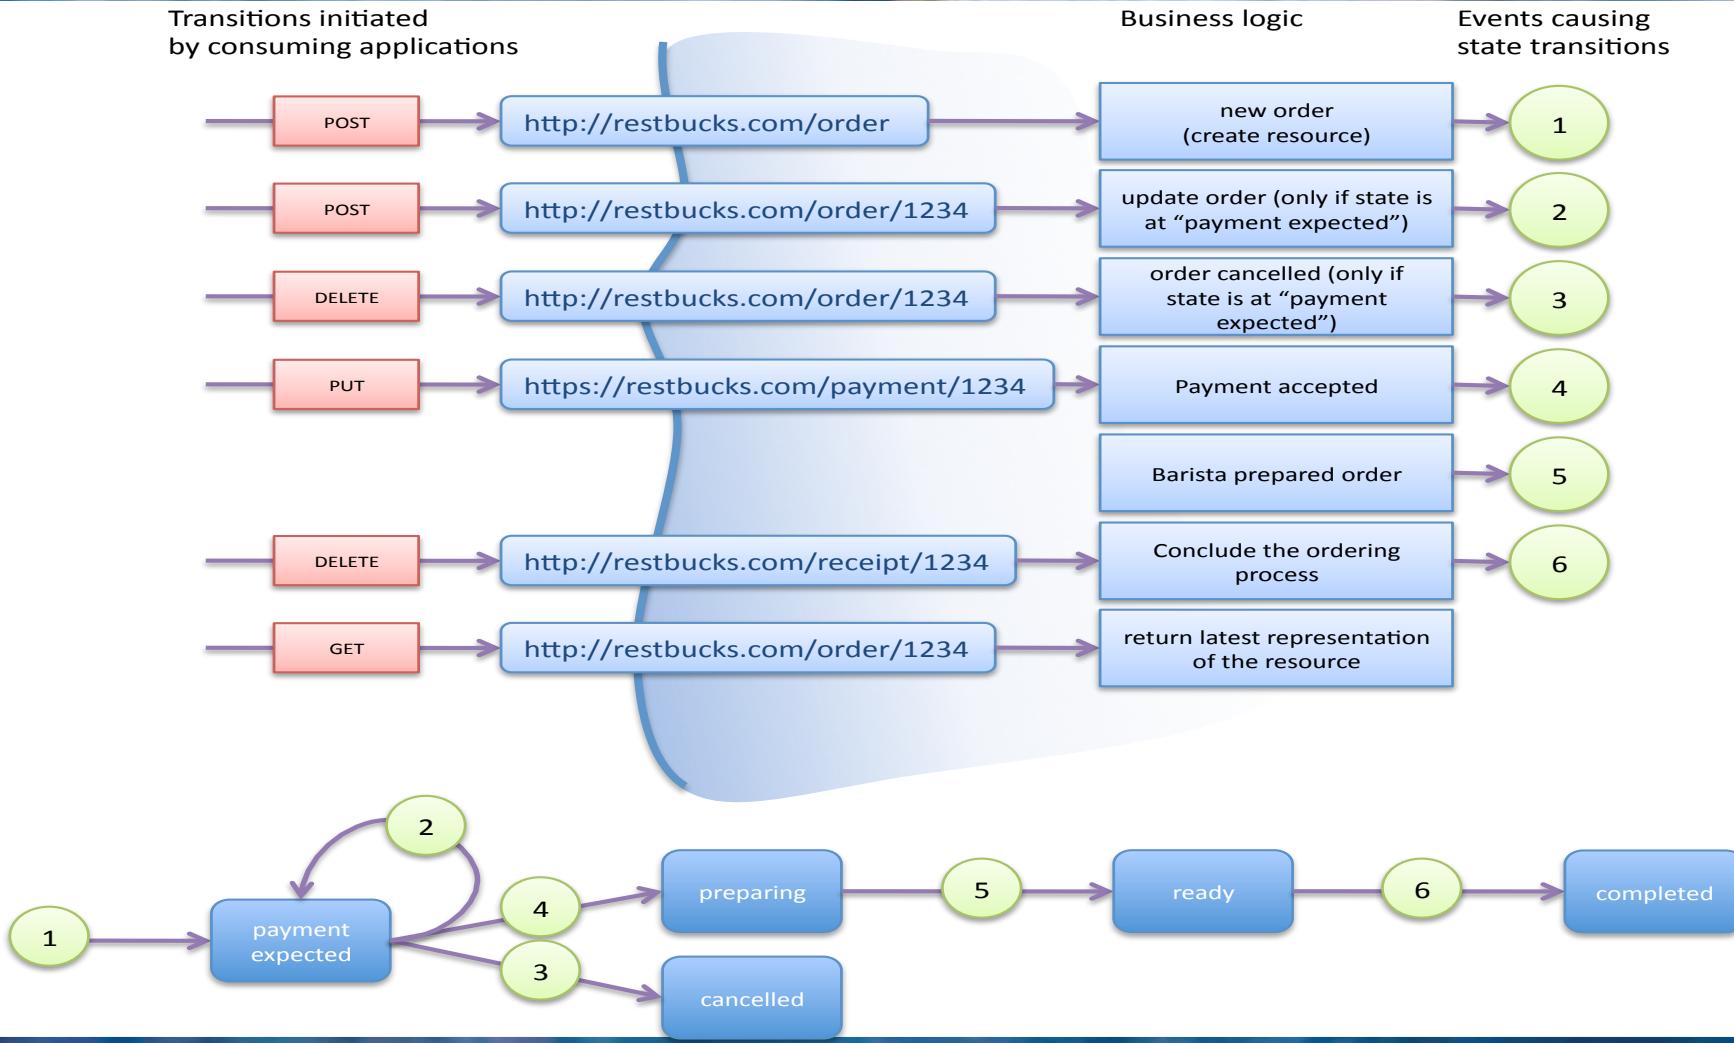
- Links declare next valid steps
- Following links and interacting with resources changes *application state*
- Media types in links define contracts
 - Media type defines processing model
 - Links (with microformats) describe state transition
- Don't need a static contract description
 - No WSDL or UML
- This is HAL JSON.
- So let's see how we order a coffee at restbucks.com...
 - Based on:
<http://www.infoq.com/articles/webber-rest-workflow>

**domain
application
protocols!**

Workflow

- How does a typical enterprise workflow look when it's implemented in a Web-friendly way?
- Let's take Restbucks ordering service an example, the happy path is:
 - Make selection
 - Add any specialities
 - Pay
 - Wait for a while
 - Collect drink

Static Interface and State Transitions



application/vnd.restbucks+xml

- What a mouthful!
- The vnd namespace is for proprietary media types
 - As opposed to the IANA-registered ones
- Restbucks own XML is a hybrid
 - We use plain old XML to convey information
 - And link elements to convey protocol
- This is important, since it allows us to create RESTful, hypermedia aware services

application/vnd.restbucks+xml



Schemas

<http://schemas.restbucks.com>



Hypermedia controls

<link> element, borrowed from Atom



Link relations

order

update

cancel (lation)

payment

receipt

self



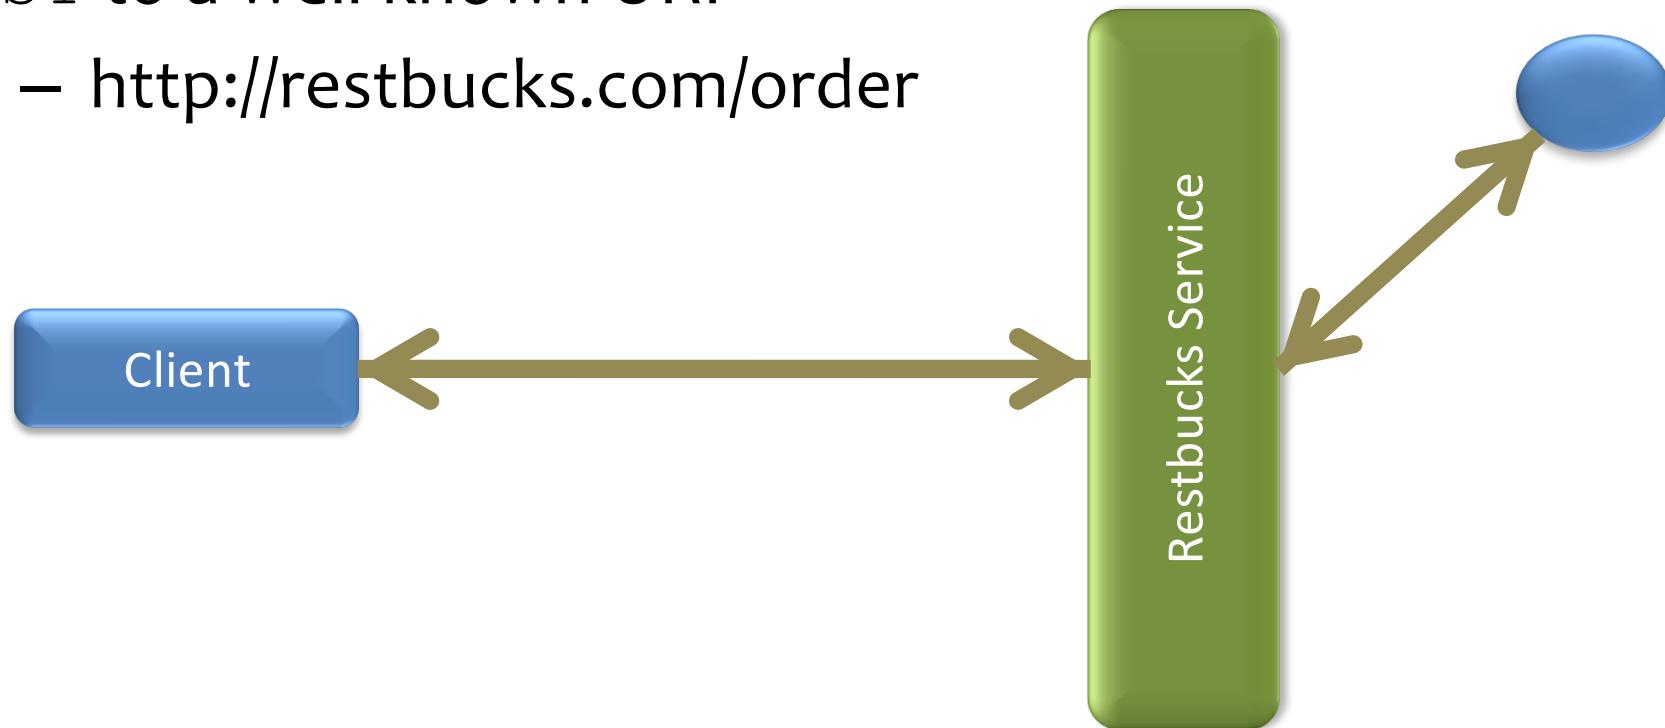
Processing model

Atom-like, links indicate HTTP idioms to apply

Place an Order

POST to a well-known URI

- <http://restbucks.com/order>



Placing an Order

Request

```
POST /order HTTP/1.1
```

```
Content-Type: application/vnd.restbucks+xml
```

```
Accept: application/vnd.restbucks+xml
```

```
Host: restbucks.com
```

```
Connection: keep-alive
```

```
Content-Length: 283
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com">
  <ns2:item>
    <ns2:milk>semi</ns2:milk>
    <ns2:size>large</ns2:size>
    <ns2:drink>latte</ns2:drink>
  </ns2:item>
  <ns2:location>takeaway</ns2:location>
</ns2:order>
```

Placing an Order

Response

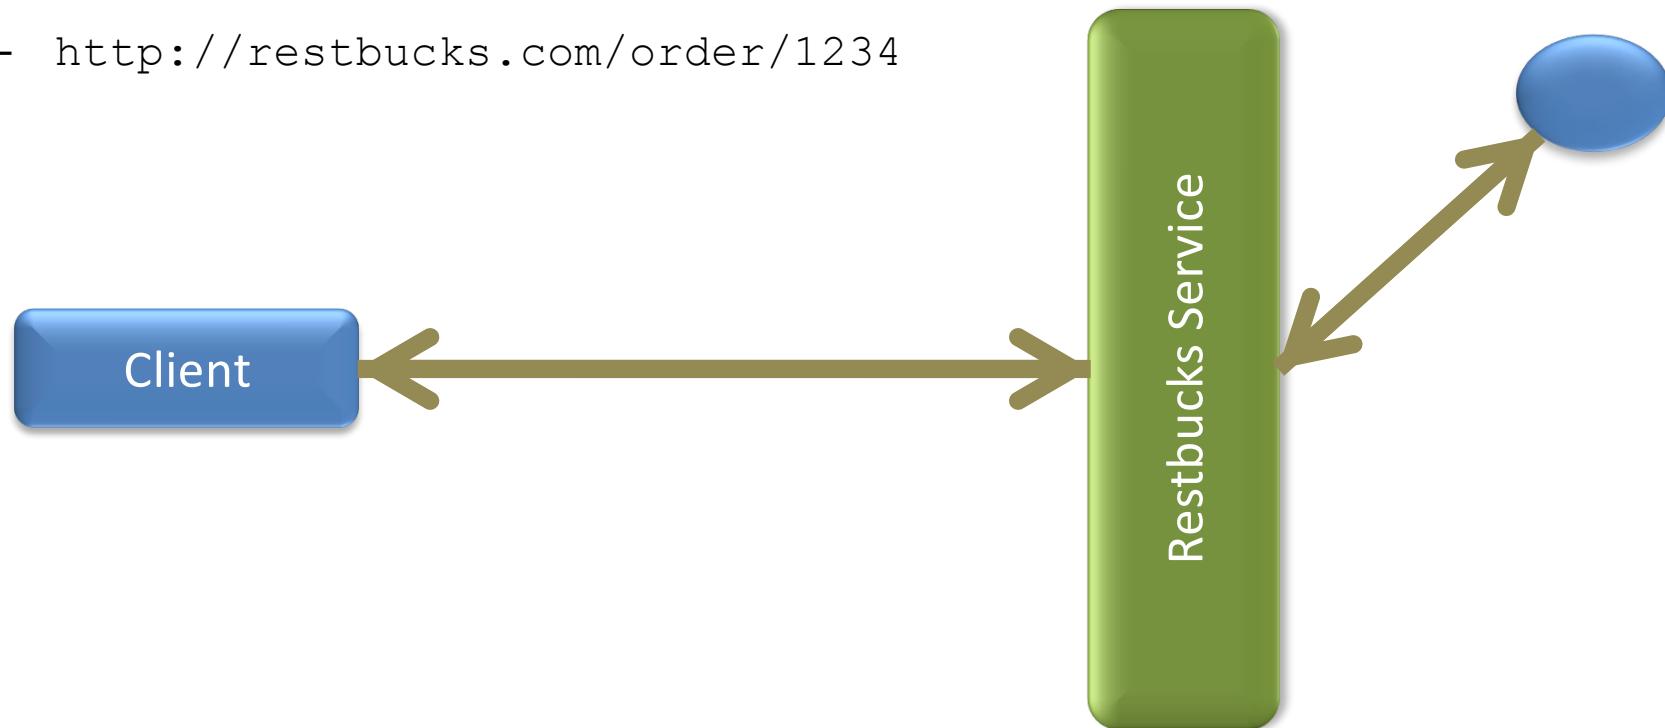
```
HTTP/1.1 201 Created
server: grizzly/1.8.1
Location: http://restbucks.com/order/1234
Content-Type: application/vnd.restbucks+xml
Content-Length: 1039
Date: Wed, 03 Mar 2010 20:58:03 GMT
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com" xmlns:ns3="http://schemas.restbucks.com/dap">
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
        rel="http://relations.restbucks.com/cancel"/>
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/payment/1234"
        rel="http://relations.restbucks.com/payment"/>
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
        rel="http://relations.restbucks.com/update"/>
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234" rel="self"/>
    <ns2:item>
        <ns2:milk>semi</ns2:milk>
        <ns2:size>large</ns2:size>
        <ns2:drink>latte</ns2:drink>
    </ns2:item>
    <ns2:location>takeaway</ns2:location>
    <ns2:cost>2.0</ns2:cost>
    <ns2:status>unpaid</ns2:status>
</ns2:order>
```

Confirm the Order

GET from the rel="self" URI

- <http://restbucks.com/order/1234>



Confirm the Order

Request

```
GET /order/1234 HTTP/1.1
```

```
Accept: application/vnd.restbucks+xml
```

```
Host: restbucks.com
```

```
Connection: keep-alive
```

Confirm the Order

Response

HTTP/1.1 200 OK

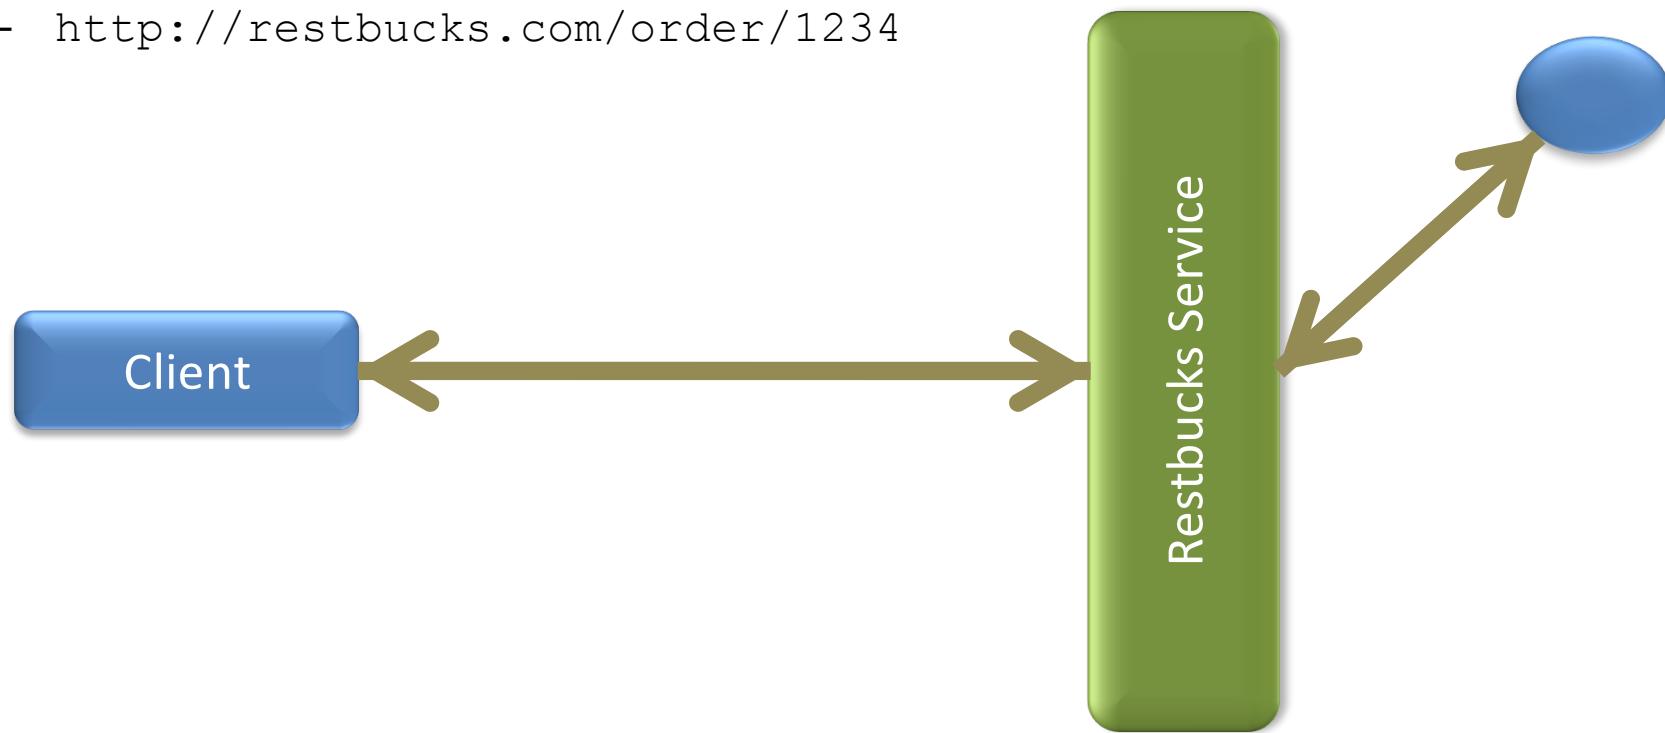
Location: http://restbucks.com/order/1234
Content-Type: application/vnd.restbucks+xml
Content-Length: 911
Date: Sun, 06 Sep 2009 06:51:22 GMT

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com" xmlns:ns3="http://schemas.restbucks.com/dap">
  <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
    rel="http://relations.restbucks.com/cancel"/>
  <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/payment/1234"
    rel="http://relations.restbucks.com/payment"/>
  <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
    rel="http://relations.restbucks.com/update"/>
  <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234" rel="self"/>
  <ns2:item>
    <ns2:milk>semi</ns2:milk>
    <ns2:size>large</ns2:size>
    <ns2:drink>latte</ns2:drink>
  </ns2:item>
  <ns2:location>takeaway</ns2:location>
  <ns2:cost>2.0</ns2:cost>
  <ns2:status>unpaid</ns2:status>
</ns2:order>
```

Change the Order

POST new order to service-generated URI

- <http://restbucks.com/order/1234>



POST? Not PUT?

- PUT expects the whole resource state to be presented in the request
 - But our clients aren't responsible for generating cost or status elements
- PATCH would be better
 - It allows us to send diffs, but is new to the family of HTTP verbs and not widely supported yet
- So POST is our only option
 - Compare this to our CRUD protocol earlier

Can we better decompose the problem space to avoid this?

Change the Order

Request

```
POST order/1234 HTTP/1.1
```

```
Content-Type: application/vnd.restbucks+xml
```

```
Accept: application/vnd.restbucks+xml
```

```
Host: restbucks.com
```

```
Connection: keep-alive
```

```
Content-Length: 288
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com">
    <ns2:item>
        <ns2:milk>semi</ns2:milk>
        <ns2:size>large</ns2:size>
        <ns2:drink>cappuccino</ns2:drink>
    </ns2:item>
    <ns2:location>takeaway</ns2:location>
</ns2:order>
```

No service-generated content

Change the Order

Response

HTTP/1.1 200 OK

Location: http://restbucks.com/order/1234
Content-Type: application/vnd.restbucks+xml
Content-Length: 916
Date: Sun, 06 Sep 2009 06:52:22 GMT

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com" xmlns:ns3="http://schemas.restbucks.com/dap">
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
        rel="http://relations.restbucks.com/cancel"/>
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/payment/1234"
        rel="http://relations.restbucks.com/payment"/>
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
        rel="http://relations.restbucks.com/update"/>
    <ns3:link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234" rel="self"/>
    <ns2:item>
        <ns2:milk>semi</ns2:milk>
        <ns2:size>large</ns2:size>
        <ns2:drink>cappuccino</ns2:drink>
    </ns2:item>
    <ns2:location>takeaway</ns2:location>
    <ns2:cost>2.0</ns2:cost>
    <ns2:status>unpaid</ns2:status>
</ns2:order>
```

Statelessness

- Remember interactions with resources are stateless
- The resource “forgets” about you while you’re not directly interacting with it
- Which means race conditions are possible
- Use `If-Unmodified-Since` on a timestamp to make sure
 - Or use `If-Match` and an ETag
- You’ll get a `412 PreconditionFailed` if you lost the race
 - But you’ll avoid potentially putting the resource into some inconsistent state

Warning: Don't be Slow!

- Can only make changes until someone actually makes your drink
 - You're safe if you use `If-Unmodified-Since` or `If-Match`
 - But resource state can change without you!

○ Request

```
PUT /order/1234 HTTP 1.1
```

```
Host: restbucks.com
```

```
...
```

○ Request

```
OPTIONS /order/1234 HTTP 1.1
```

```
Host: restbucks.com
```

○ Response

```
409 Conflict
```

○ Response

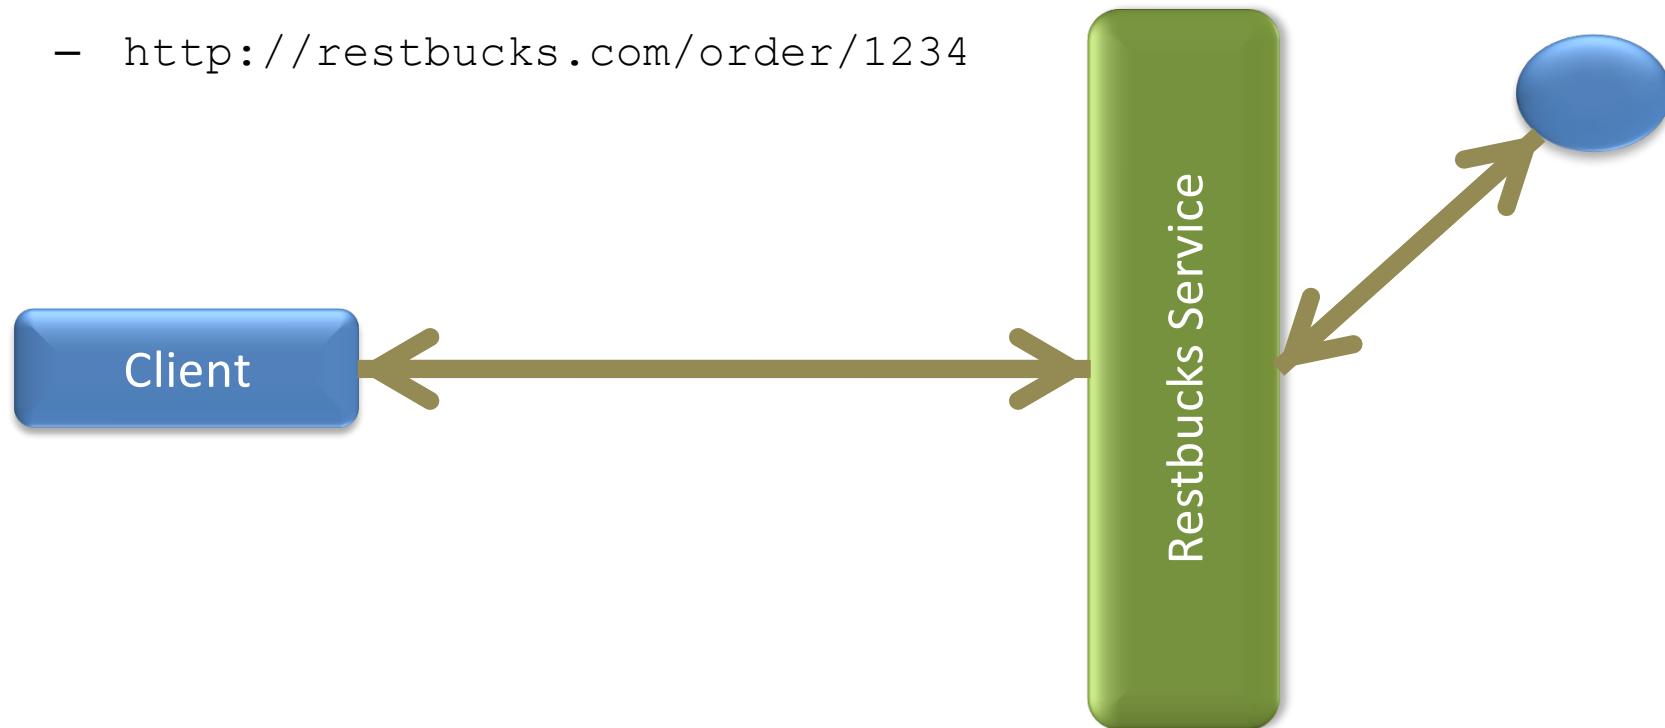
```
Allow: GET
```

Too slow!
Someone
else has
changed the
state of my
order

What if we want to cancel?

DELETE order at rel="self" URI

- <http://restbucks.com/order/1234>



What if we want to cancel?

Request

```
DELETE /order/1234 HTTP/1.1
```

```
Host: restbucks.com
```

```
Connection: keep-alive
```

Response

```
HTTP/1.1 200 OK
```

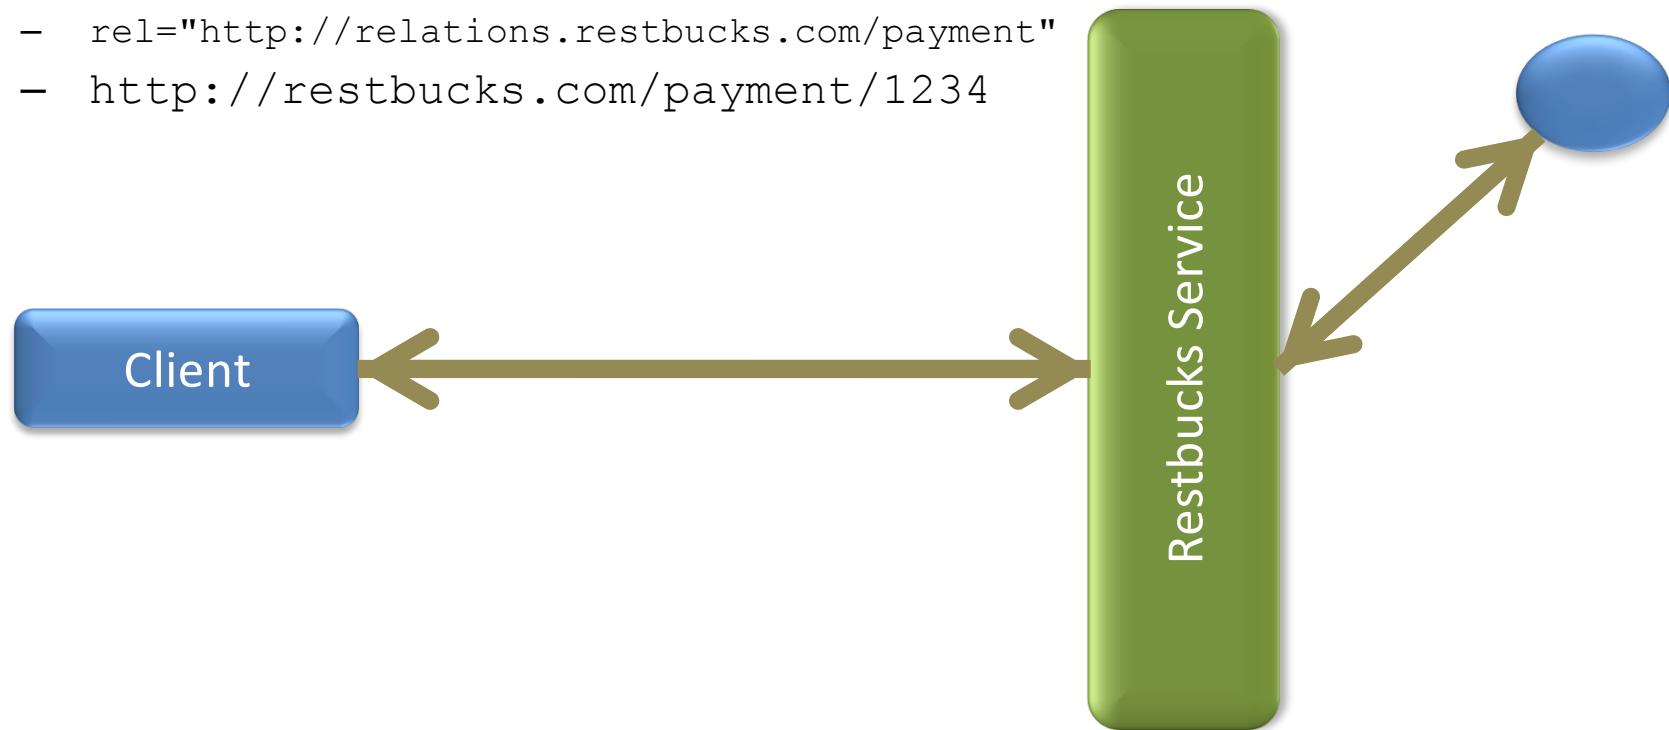
```
Content-Type: application/vnd.restbucks+xml
```

```
Date: Sun, 06 Sep 2009 06:53:22 GMT
```

Order Payment

PUT to service-generated payment URI

- rel="http://relations.restbucks.com/payment"
- http://restbucks.com/payment/1234



Why PUT this time?

- It's idempotent
 - Safe to repeat in the event of failures
- For all \$ transactions, prefer idempotent verbs

Order Payment

Request

```
PUT /payment/1234 HTTP/1.1
```

```
Content-Type: application/vnd.restbucks+xml
```

```
Accept: application/vnd.restbucks+xml
```

```
Host: restbucks.com
```

```
Connection: keep-alive
```

```
Content-Length: 287
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<payment xmlns="http://schemas.restbucks.com">
    <amount>2.0</amount>
    <cardholderName>Michael Farraday</cardholderName>
    <cardNumber>11223344</cardNumber>
    <expiryMonth>12</expiryMonth>
    <expiryYear>12</expiryYear>
</payment>
```

Order Payment

Response

HTTP/1.1 201 Created

server: grizzly/1.8.1

Location: http://restbucks.com/payment/1234

Content-Type: application/vnd.restbucks+xml

Content-Length: 650

Date: Wed, 03 Mar 2010 20:58:03 GMT

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:payment xmlns="http://schemas.restbucks.com/dap" xmlns:ns2="http://schemas.restbucks.com">
  <link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/order/1234"
    rel="http://relations.restbucks.com/order"/>
  <link mediaType="application/vnd.restbucks+xml" uri="http://restbucks.com/receipt/1234"
    rel="http://relations.restbucks.com/receipt"/>
  <ns2:amount>2.0</ns2:amount>
  <ns2:cardholderName>Michael Farraday</ns2:cardholderName>
  <ns2:cardNumber>11223344</ns2:cardNumber>
  <ns2:expiryMonth>12</ns2:expiryMonth>
  <ns2:expiryYear>12</ns2:expiryYear>
</ns2:payment>
```

What if we want to cancel now?

Request

```
DELETE /order/1234 HTTP/1.1
```

```
Host: restbucks.com
```

```
Connection: keep-alive
```

Response

```
HTTP/1.1 405 Method Not Allowed
```

```
Allow: GET
```

```
Content-Type: application/vnd.restbucks+xml
```

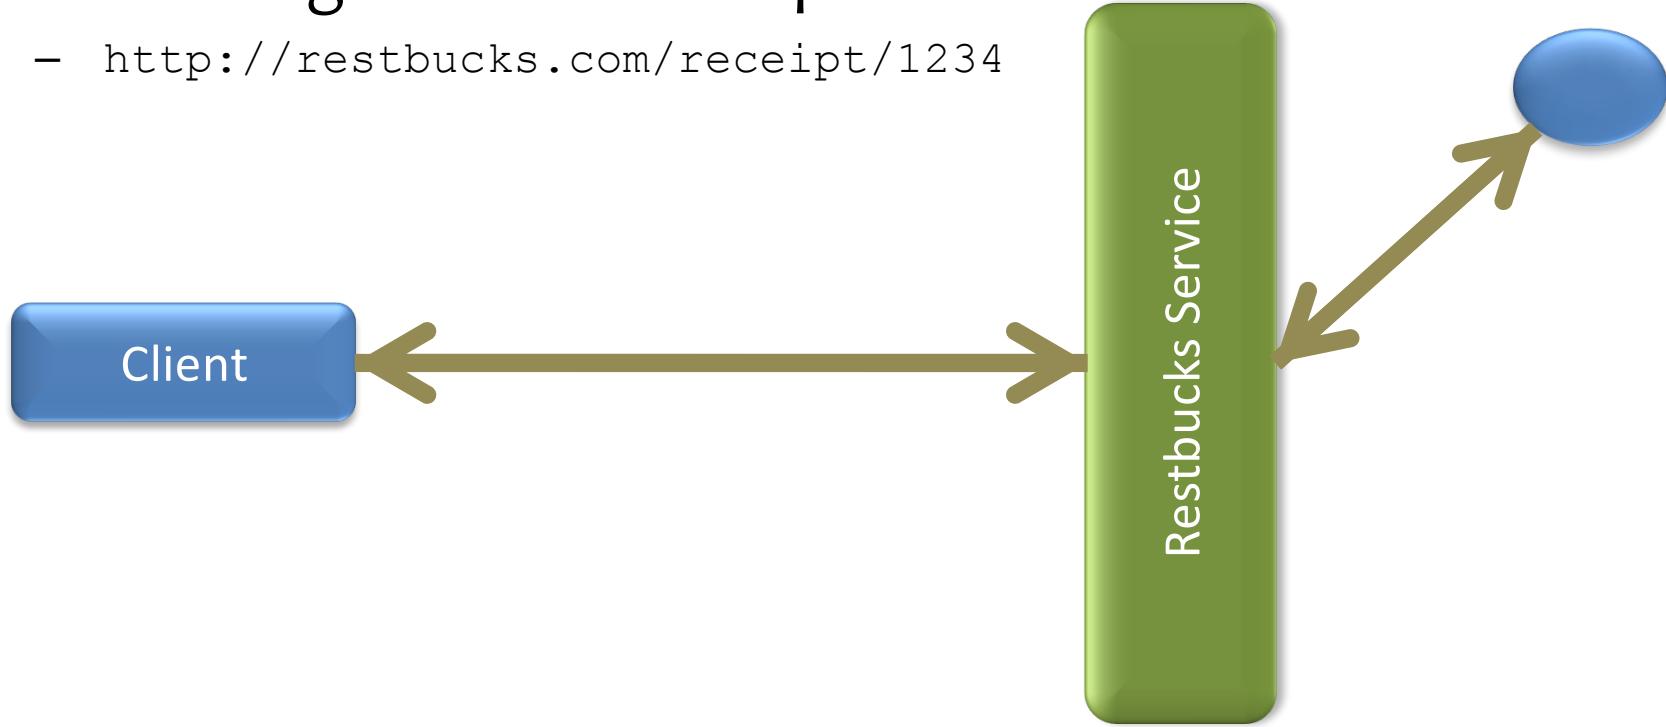
```
Content-Length: 0
```

```
Date: Sun, 06 Sep 2009 07:43:29 GMT
```

Grab a Receipt

GET service-generated receipt URI

- <http://restbucks.com/receipt/1234>



Grab a Receipt

Request

GET /receipt/1234 HTTP/1.1

Accept: application/vnd.restbucks+xml

Host: restbucks.com

Grab a Receipt

- Request

HTTP/1.1 200 OK

server: grizzly/1.8.1

Content-Type: application/vnd.restbucks+xml

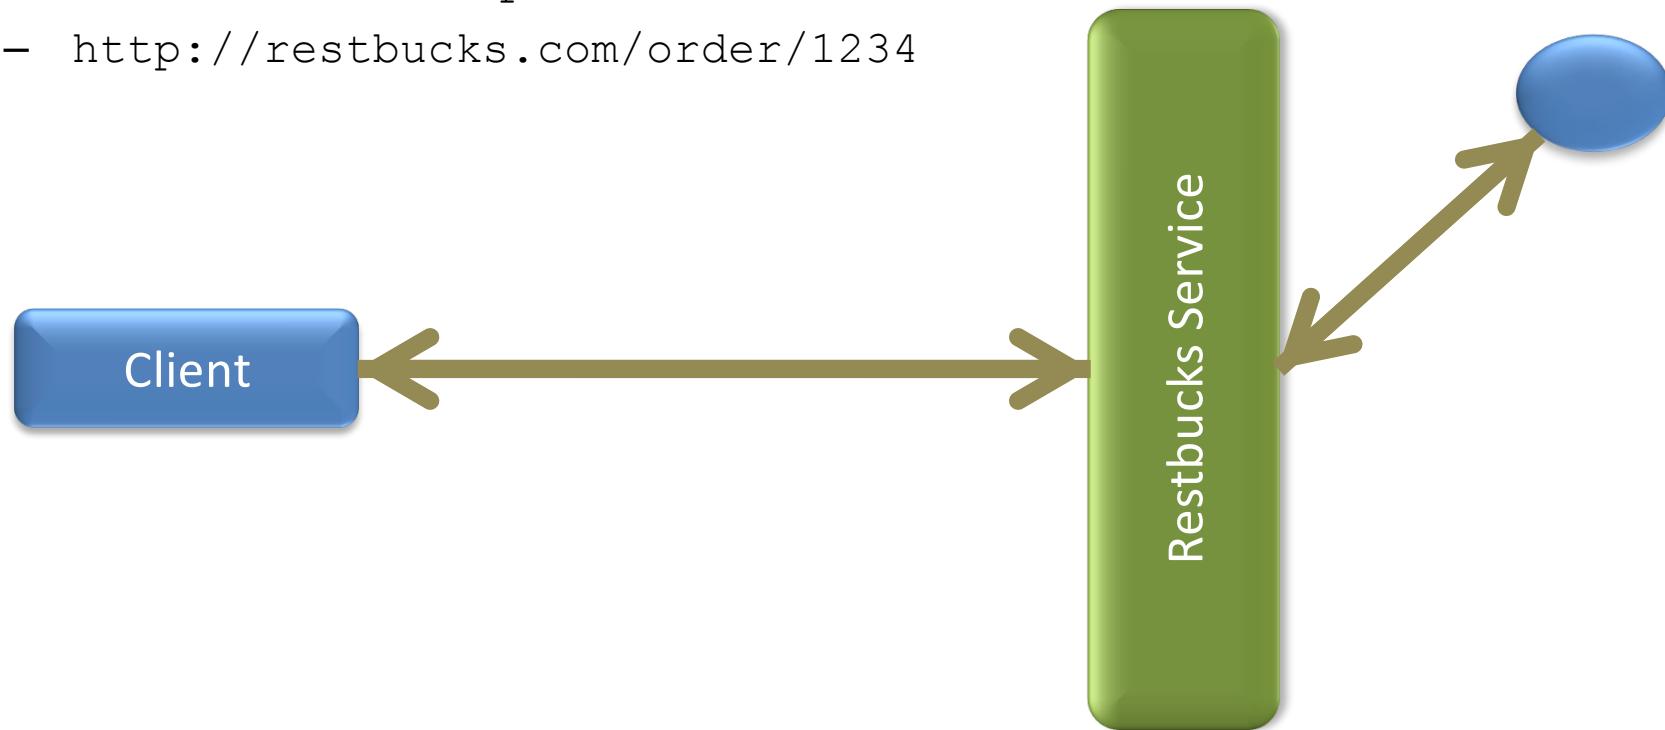
Content-Length: 384

Date: Wed, 03 Mar 2010 20:58:03 GMT

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:receipt xmlns="http://schemas.restbucks.com/dap"
  xmlns:ns2="http://schemas.restbucks.com">
  <link mediaType="application/vnd.restbucks+xml"
    uri="http://restbucks.com/order/1234"
    rel="http://relations.restbucks.com/order"/>
  <ns2:amount>2.0</ns2:amount>
  <ns2:paid>2010-03-03T21:58:03.834+01:00</ns2:paid>
</ns2:receipt>
```

Poll until the order's ready

- GET from `rel="http://relations.restbucks.com/order"` URI
 - `http://restbucks.com/order/1234`



Polling for a ready order on the wire

Request

GET /order/1234 HTTP/1.1

Host: restbucks.com

Connection: keep-alive

Order's ready, take the receipt and walk away

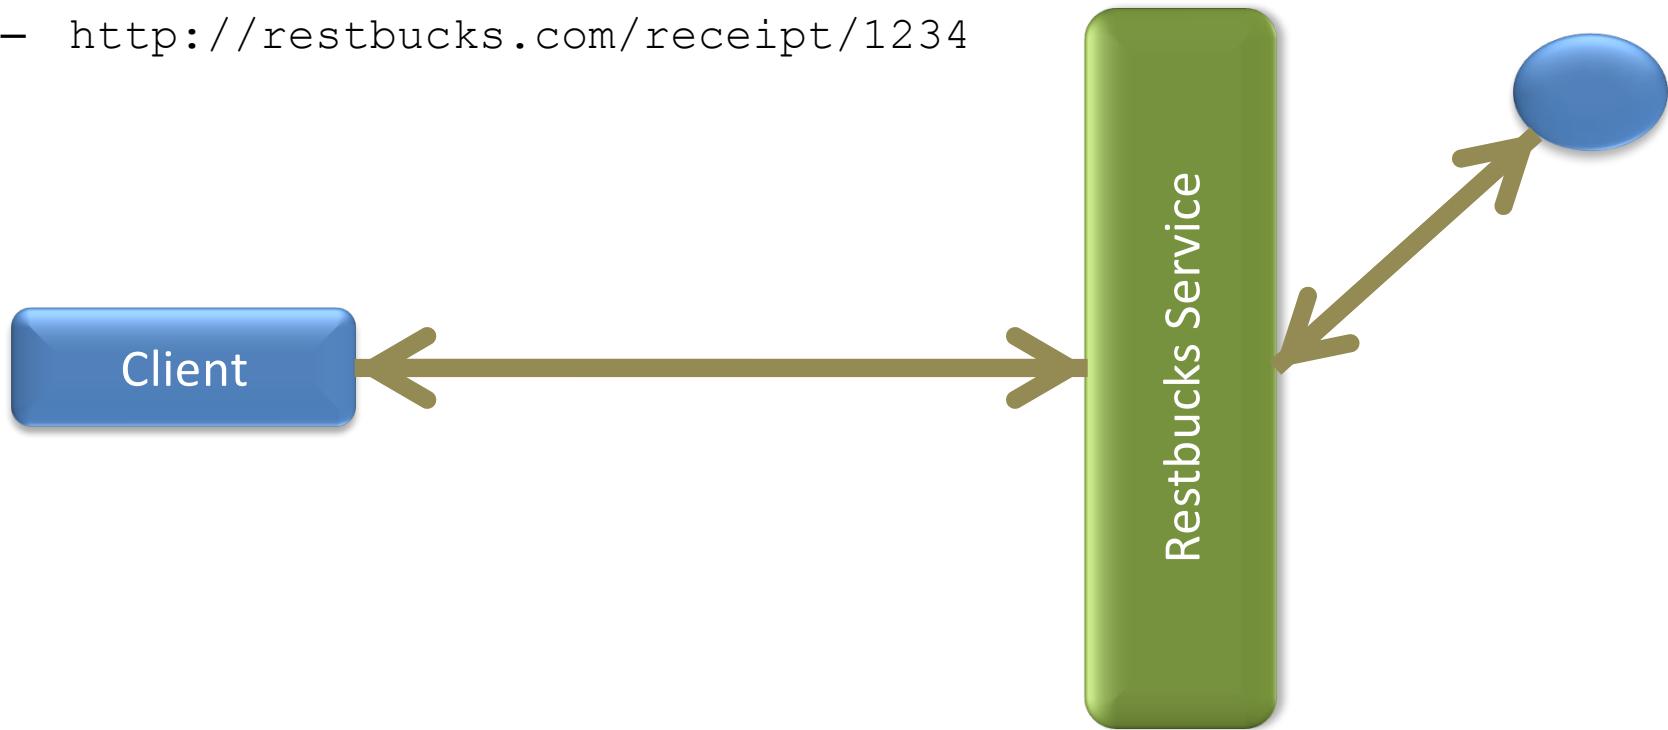
Response

```
HTTP/1.1 200 OK
server: grizzly/1.8.1
Content-Type: application/vnd.restbucks+xml
Content-Length: 534
Date: Wed, 03 Mar 2010 20:58:03 GMT
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com"
    xmlns:ns3="http://schemas.restbucks.com/dap">
    <ns3:link mediaType="application/vnd.restbucks+xml"
        uri="http://restbucks.com/receipt/1234"
        rel="http://relations.restbucks.com/reciept"/>
    <ns2:item>
        <ns2:milk>semi</ns2:milk>
        <ns2:size>large</ns2:size>
        <ns2:drink>cappuccino</ns2:drink>
    </ns2:item>
    <ns2:location>takeaway</ns2:location>
    <ns2:cost>2.0</ns2:cost>
    <ns2:status>ready</ns2:status>
</ns2:order>
```

Complete the Order

DELETE order at rel="http://relations.restbucks.com/receipt" URI
– <http://restbucks.com/receipt/1234>



Complete the Order

Request

DELETE/receipt/1234 HTTP/1.1

Host: restbucks.com

Connection: keep-alive

Complete the Order

Response

```
HTTP/1.1 200 OK
server: grizzly/1.8.1
Content-Type: application/xml
Content-Length: 393
Date: Wed, 03 Mar 2010 20:58:03 GMT
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns2:order xmlns:ns2="http://schemas.restbucks.com"
             xmlns:ns3="http://schemas.restbucks.com/dap">
    <ns2:item>
        <ns2:milk>semi</ns2:milk>
        <ns2:size>large</ns2:size>
        <ns2:drink>cappuccino</ns2:drink>
    </ns2:item>
    <ns2:location>takeaway</ns2:location>
    <ns2:cost>2.0</ns2:cost>
    <ns2:status>taken</ns2:status>
</ns2:order>
```

No Hypermedia
Controls

Finally drink your coffee...



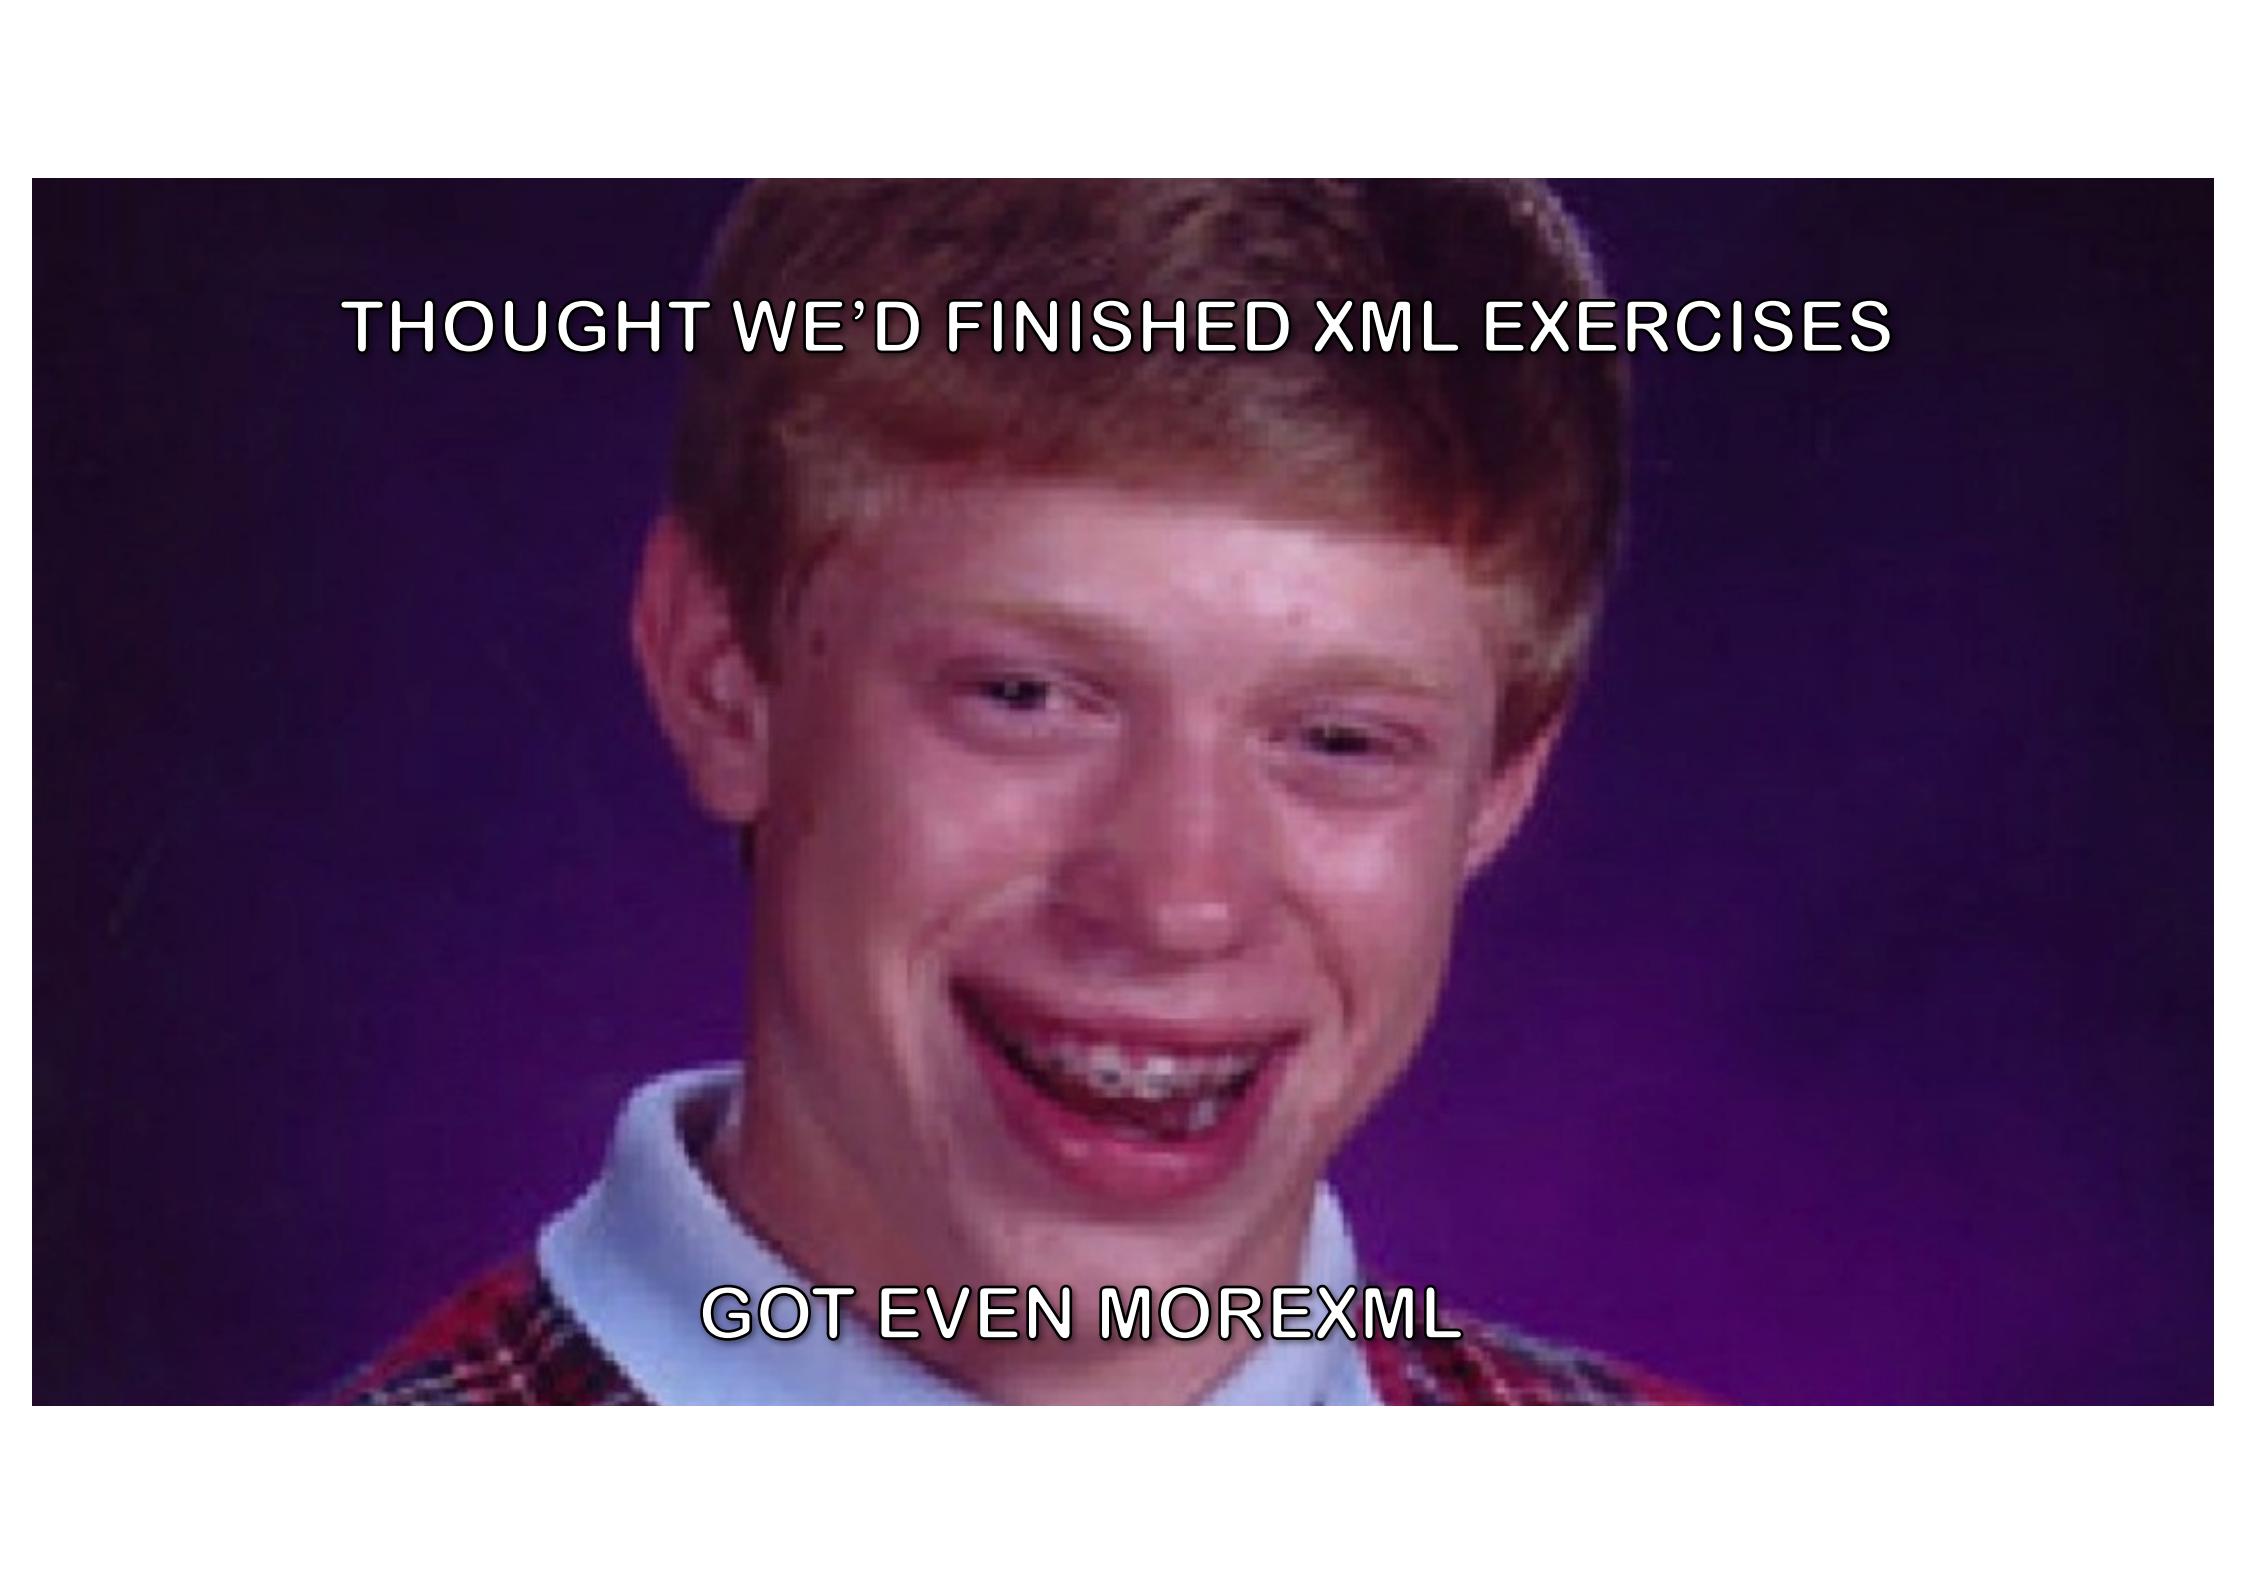
Source: http://images.businessweek.com/ss/06/07/top_brands/image/starbucks.jpg

Ordering and Payments Protocol Weaknesses

- Lacks HTTP idioms
 - ETags
 - Cache metadata
 - Etc
- No circuit breaker
 - E.g. 413 at busy times

Exercise

- How on Earth did you pay for a coffee in the CRUD exercise?
- Some of you picked out URIs from the response entity
- But that's a job that can be done by a computer, so...
- Let's properly order a coffee starting at /order



THOUGHT WE'D FINISHED XML EXERCISES

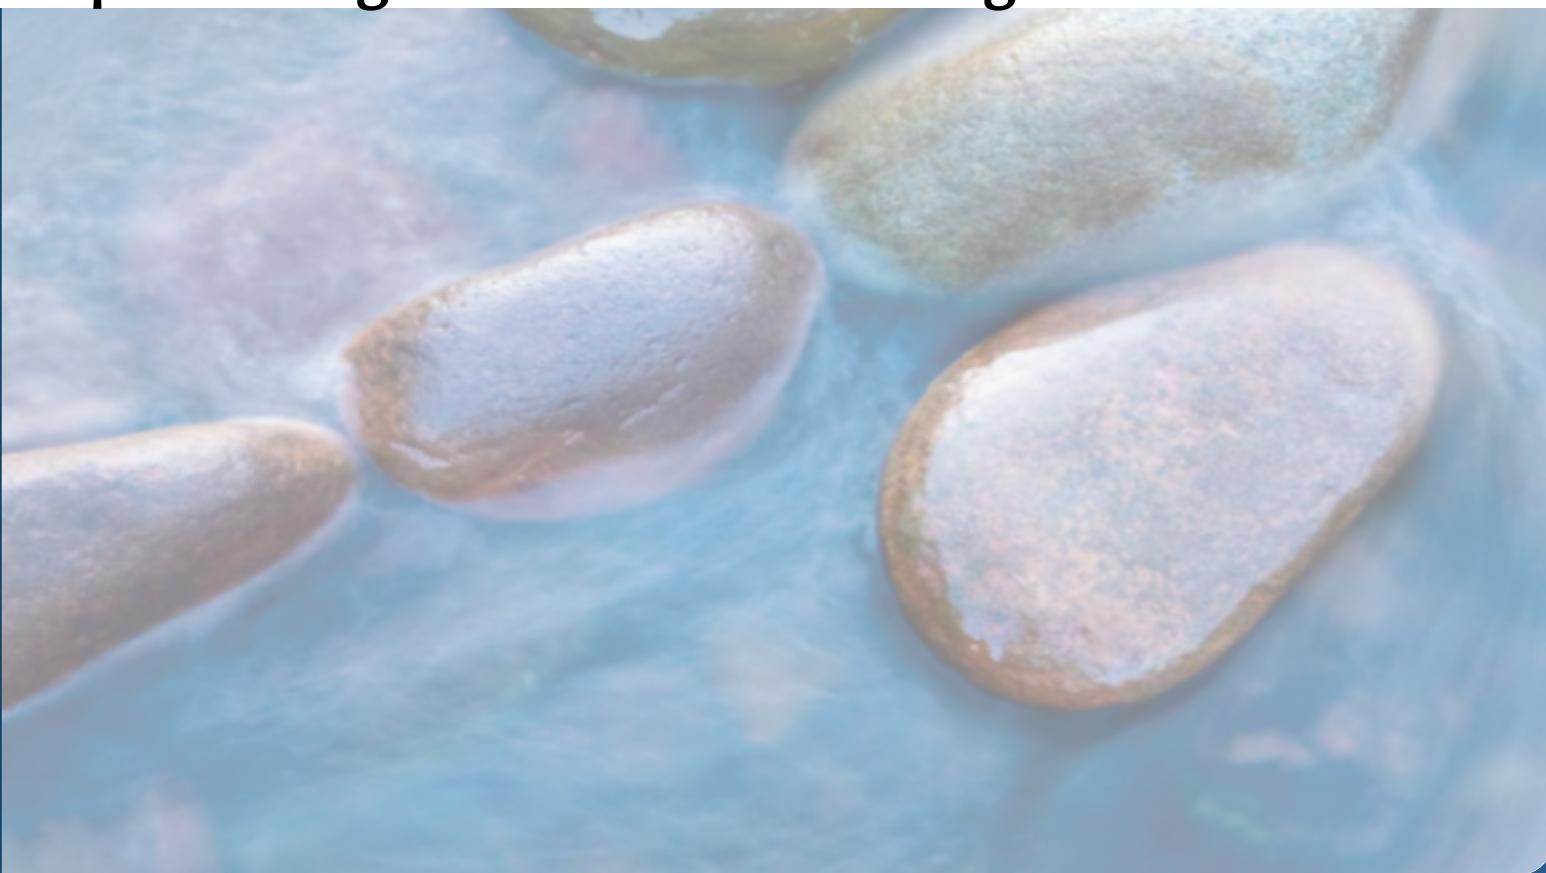
GOT EVEN MOREXML

Domain Application Protocol

- The XML contains links which describe* a protocol
- If you bind to those link names, then the URIs don't matter
- Treat them as “things you'd like to do next”
- Let's code/curl

Tech Interlude

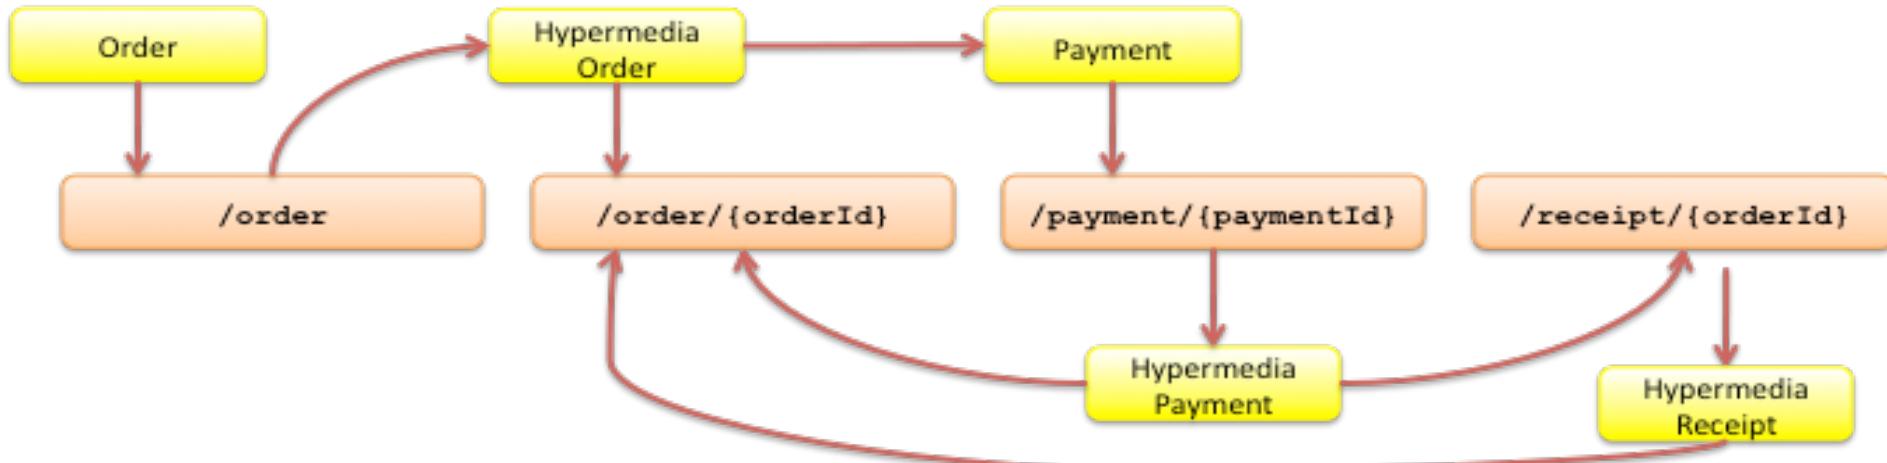
Implementing the Restbucks Ordering Service



Java

patterns subsequently captured by Restfulie

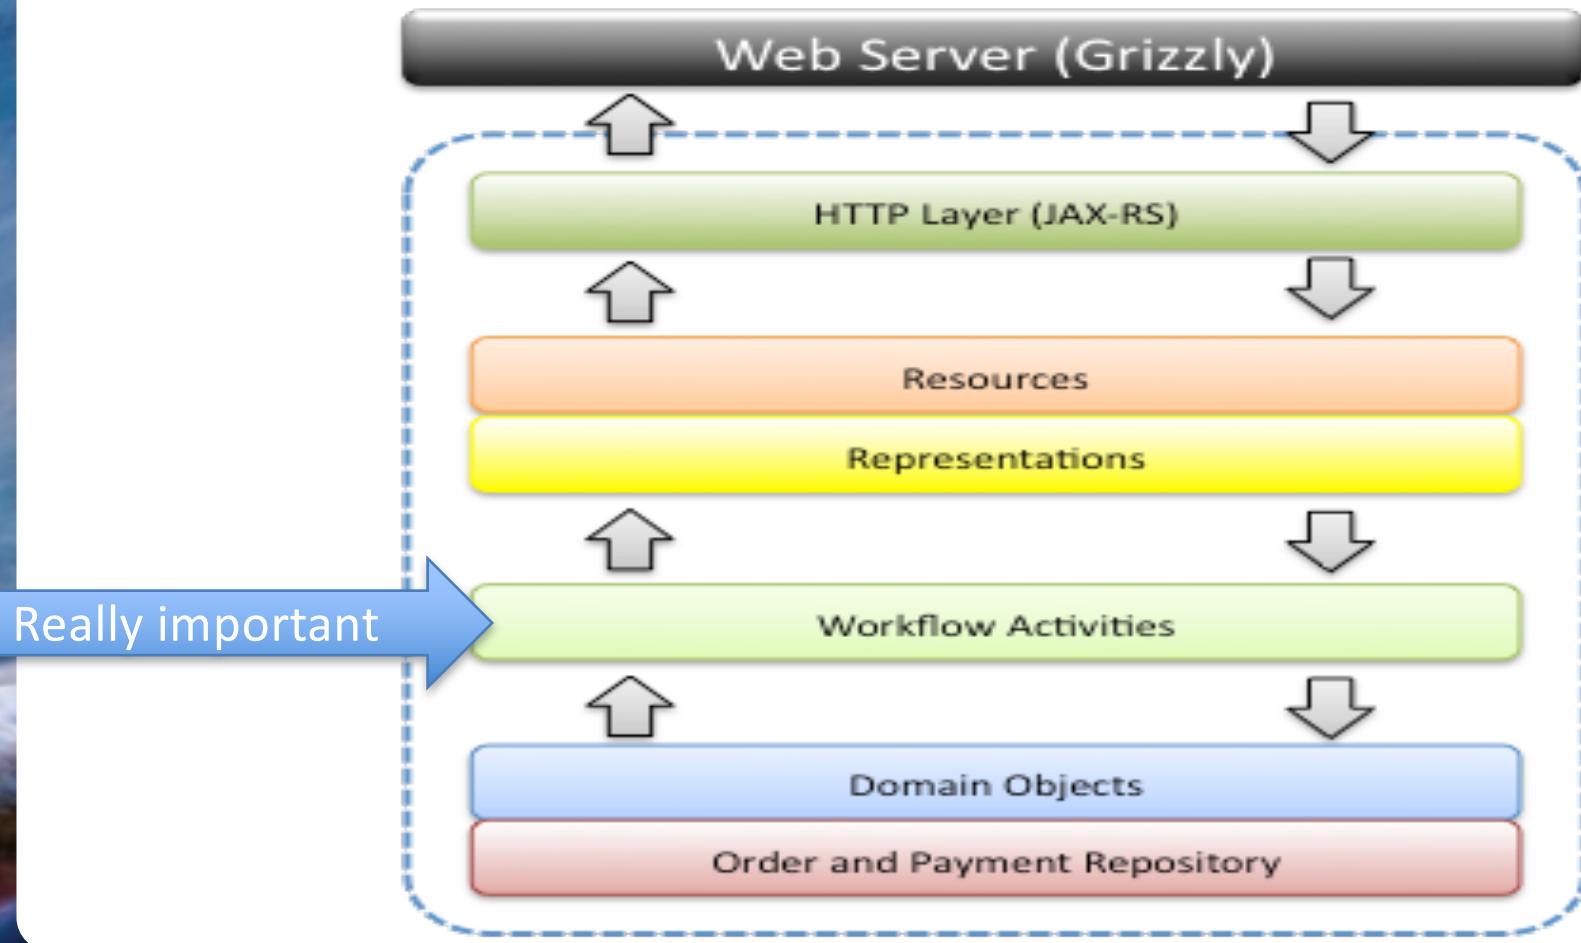
Protocol/Resource Mapping



Server's easy

- We picked Jersey (JAX-RS)
- Somewhat able to TDD it, with effort
 - Some of the server-side return types don't invite easy inspection sadly

Java Server-side Architecture



Typical JAX-RS Plumbing

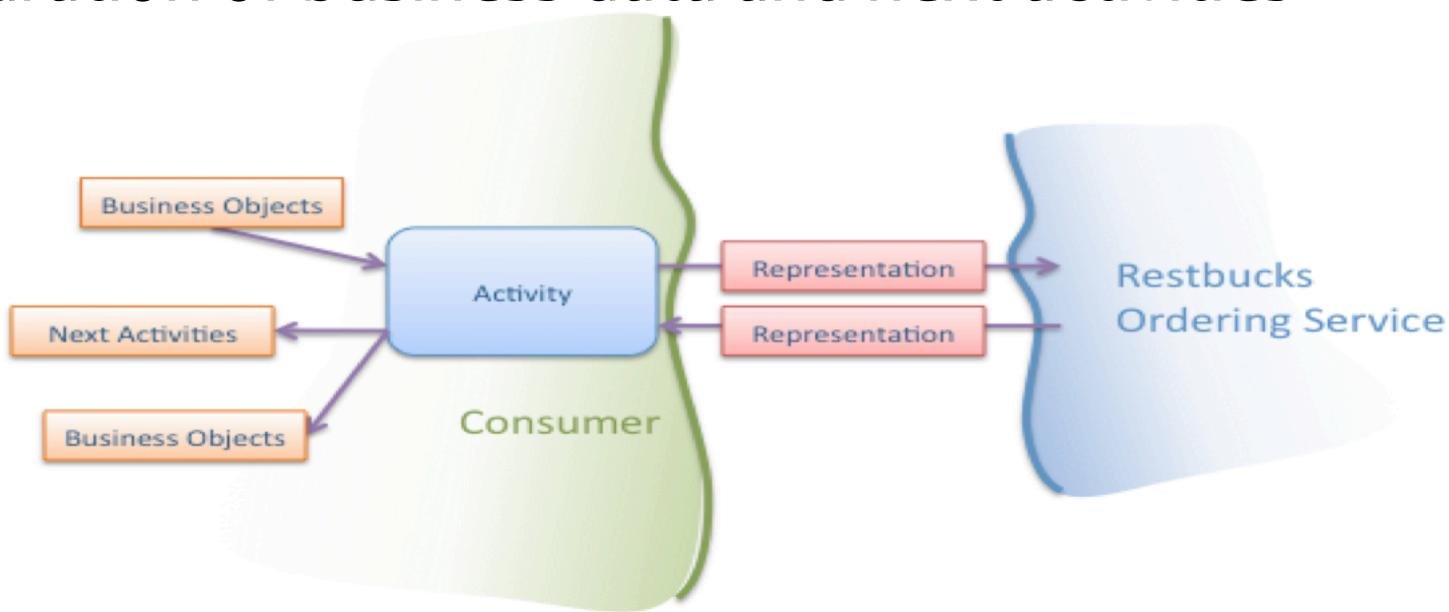
```
@POST  
@Path("/{orderId}")  
@Consumes("application/vnd.restbucks+xml")  
@Produces("application/vnd.restbucks+xml")  
public Response updateOrder(String orderRepresentation) {  
    try {  
        OrderRepresentation responseRepresentation = new UpdateOrderActivity()  
            .update(  
                OrderRepresentation.fromXmlString(  
                    orderRepresentation)  
            .getOrder(), new RestbucksUri(uriInfo.getRequestUri()));  
        return Response.ok().entity(responseRepresentation).build();  
    } catch (InvalidOrderException ioe) {  
        return Response.status(Status.BAD_REQUEST).build();  
    } catch (NoSuchOrderException nsoe) {  
        return Response.status(Status.NOT_FOUND).build();  
    } catch(UpdateException ue) {  
        return Response.status(Status.CONFLICT).build();  
    } catch (Exception ex) {  
        return Response.serverError().build();  
    }  
}
```

Generating a Response Representation

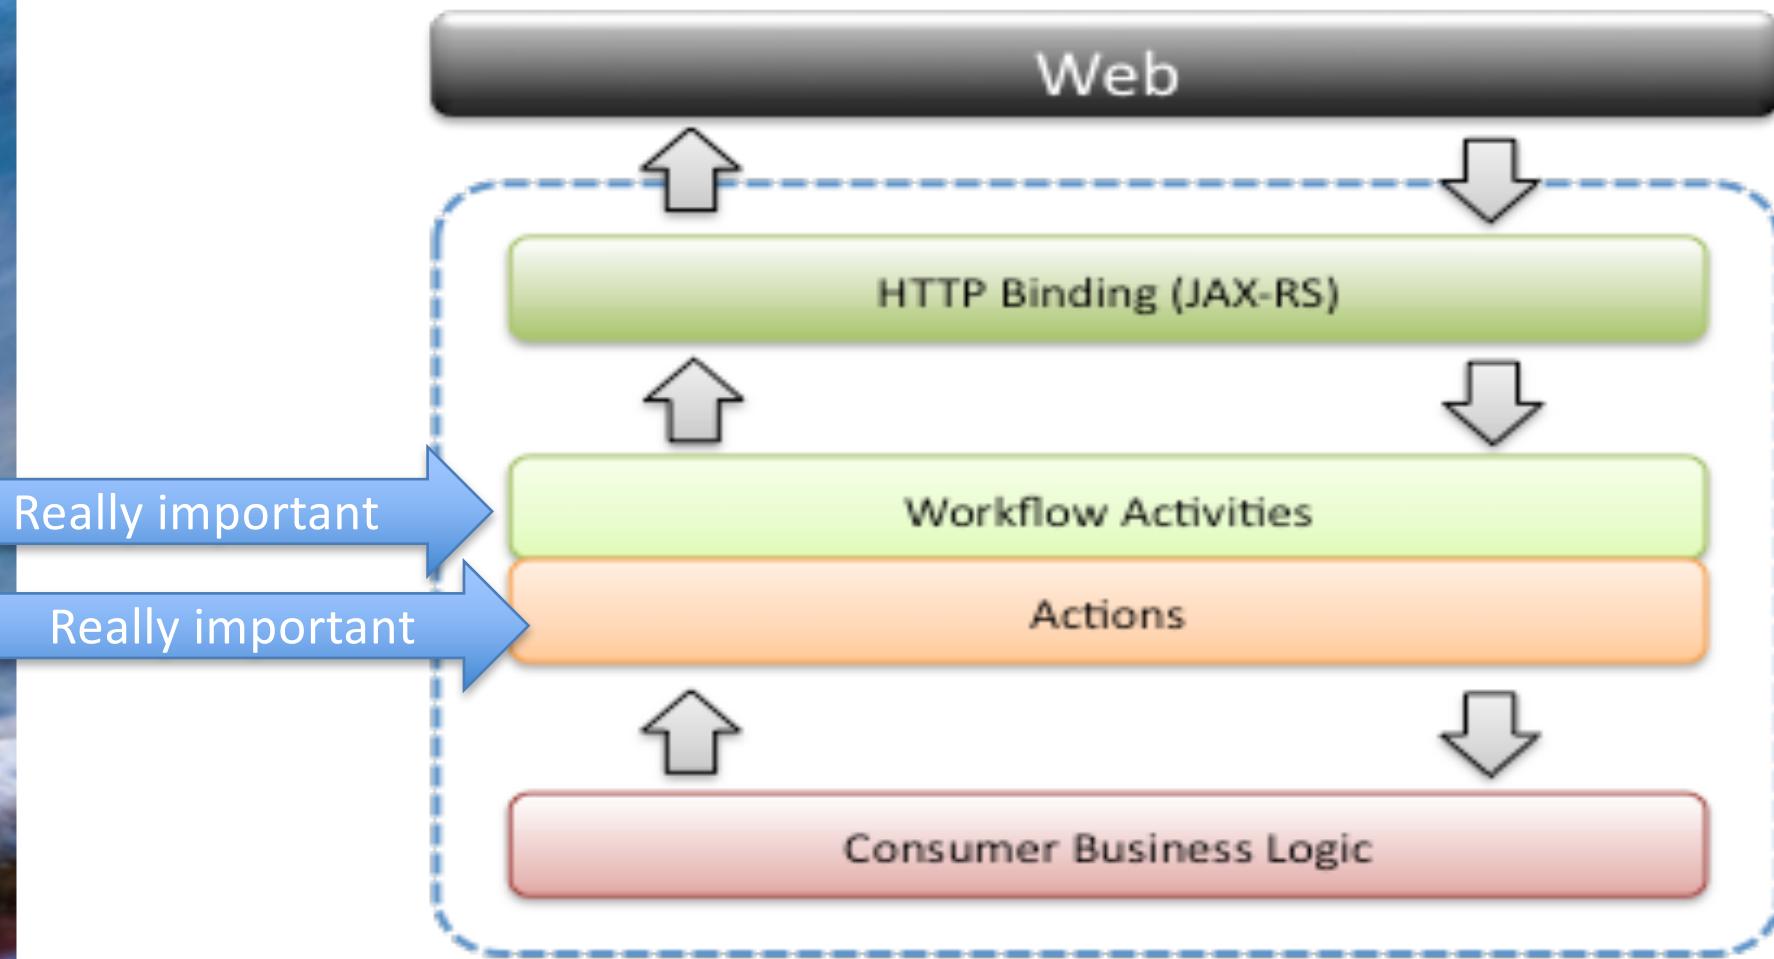
```
public class CreateOrderActivity {  
    public OrderRepresentation create(Order order,  
                                     RestbucksUri requestUri) {  
        order.setStatus(OrderStatus.UNPAID);  
  
        Identifier identifier = OrderRepository.current().store(order);  
  
        RestbucksUri orderUri = new RestbucksUri(requestUri.getBaseUri() +  
                                                 "/order/" + identifier.toString());  
  
        RestbucksUri paymentUri = new RestbucksUri(requestUri.getBaseUri() +  
                                                 "/payment/" + identifier.toString());  
        return new OrderRepresentation(order,  
                                      new Link(Representation.RELATIONS_URI + "cancel", orderUri),  
                                      new Link(Representation.RELATIONS_URI + "payment", paymentUri),  
                                      new Link(Representation.RELATIONS_URI + "update", orderUri),  
                                      new Link(Representation.SELF_REL_VALUE, orderUri));  
    }  
}
```

Consumers are harder

- We chose Jersey client
- Code to the media type, be declarative
- Separation of business data and next activities



Java Consumer-side Architecture



Consumer-side workflow

```
public void orderAndPay(Order order, URI entryPointUri) {  
  
    PlaceOrderActivity placeOrderActivity = new PlaceOrderActivity();  
    placeOrderActivity.placeOrder(order, entryPointUri);  
  
    // Order processing ommitted for brevity...  
  
    Actions actions = placeOrderActivity.getActions();  
  
    if(actions.has(PaymentActivity.class)) {  
        PaymentActivity paymentActivity = actions.get(PaymentActivity.class);  
        paymentActivity.payForOrder(  
            payment().withAmount(readOrderActivity.getOrder().getCost())  
                .build());  
        actions = paymentActivity.getActions();  
    }  
    // Remainder of workflow ommitted for brevity  
}
```

Summary

- Web basics
- HTTP fundamentals
- Web-as-platform
- URI templates are not a good idea
- Crud is good
- Hypermedia is better!
- Getting a coffee